

Asad Ikram

asad.ikram53@gmail.com · +44 7482 376417 · Lahore, Pakistan · linkedin.com/in/asad-ikram98 · asad-ikram-mc.github.io/portfolio

Professor Øystein Sæbø, Head of Department
Department of Information Systems, Faculty of Social Sciences
University of Agder, Kristiansand · September 2026

Re: PhD Research Fellow in Information Systems on AI, Digital Platforms and Extremism (Jobbnorge 307398)

Dear Professor Sæbø,

I am applying for the PhD Research Fellowship on AI, digital platforms and extremism. My proposal, "Counting What the Classifier Sees", is enclosed; this letter is about why I want to do that work here.

The call asks how platforms circulate radical content, how algorithms amplify or suppress it, and how that can be governed. I have spent seven years building the systems that produce the numbers those questions are answered with. At M+C Saatchi Fluency I run a pipeline that ingests ten platforms daily; at Dubizzle Labs I led a team running over five hundred collectors across fifteen countries; at ArtemisAI, which I founded, we labelled 3.39 million comments with a two-tier annotation system. Looking honestly at that last system changed what I wanted to do. Our two tiers agree on 82.9 percent of sentiment labels and 0.0 percent of toxicity labels, while the toxicity model reports 96.7 percent confidence. The dashboard called that accuracy. It was inter-model agreement with nothing behind it, which is what happens when instruments are adopted faster than the discipline to check them.

Every prevalence and amplification figure in the extremism literature is produced the same way, by a classifier whose agreement with other classifiers is rarely carried into the final number. My methods paper with Dr Ella Haig shows, in a five-hundred-replication simulation, that pooling platform scores at face value gets the sign of the aggregate wrong on about a quarter of days. The proposal applies that machinery to extremist content across platforms and across English and Norwegian, re-estimates amplification with the error carried through, and asks, through the affordance-based institutional-logics lens of your recent Information Systems Journal work, how platforms, regulators and NIVE each decide what counts as evidence. It ends with an artefact: a validity protocol a national centre can use, which matters because NIVE is mandated to build a common risk assessment tool in a family the latest Campbell review finds has no demonstrated predictive validity.

I should be plain about what I do not bring. I have no background in extremism studies, no Norwegian, and no research experience of qualitative interviewing. The proposal is designed around those gaps rather than over them: the codebook is co-developed with practitioners, the language gap becomes a cross-language invariance test, and the interview study is sized and supervised accordingly. I will start learning Norwegian on arrival.

On eligibility: my four-year BSc in Computer Science and one-year MSc in Data Analytics (Distinction, masters project 78) give five years of education. HK-dir's automatic recognition equates the UK master's to a Norwegian master's of 1.5 years; I have applied for ordinary recognition of the bachelor's and can supply the case number. Ms Færavaag kindly confirmed on 4 September that I may apply on this basis.

I would be glad to discuss any of it.

Yours sincerely,
Asad Ikram