

Asad Ikram

MSc Data Analytics (Distinction), University of Portsmouth · Chevening Scholar 2024/25
Applicant, Research Assistant / PhD Student, Geography of Code group, Corvinus University of Budapest
Lahore, Pakistan · +44 7482 376417 · asad.ikram53@gmail.com · LinkedIn: [linkedin.com/in/asad-ikram98](https://www.linkedin.com/in/asad-ikram98) · GitHub: github.com/asadfix

PROFILE

Data engineer and researcher; my own career is a case of platform-mediated software work. I started freelancing on Fiverr in 2018 as a computer science student in Pakistan (300+ completed orders), then built a data career at Dubizzle Labs and in remote work for employers and clients in Canada, Italy, Ireland, the United Arab Emirates, the United States and the United Kingdom. I work daily in Python and SQL on large, messy, multilingual web and social media data, build and evaluate machine-learning and LLM labelling systems, and build AI agents into data systems (a self-healing crawler agent and a crawler-builder agent), with AI coding assistants supporting my own work. I want to study how platforms and AI change software work with GitHub, LinkedIn and other trace data, and what those traces can and cannot measure, applying what I learned in the MSc, at Fluency and at Artemis AI, and carrying the research back into how social media and data products are built.

EDUCATION

Sep 2024 - Sep 2025 awarded 8 Dec 2025	MSc Data Analytics, Distinction. University of Portsmouth, United Kingdom. Chevening Scholar. Credit-weighted average 70.5/100; master's project 78/100. Supervisor: Dr Ella Haig. Modules included Big Data and Machine Learning, Intelligent Data and Text Analytics, Explainable AI and Data Management. Taught and assessed in English. Dissertation: <i>Cross-Platform Sentiment Analysis of Public Reaction to UK Economic Policies</i> . About 279,000 public posts from X/Twitter, Reddit, YouTube and Quora, under university ethics approval (TETHIC-2025-111094). Fine-tuned BERT against BiLSTM and lexicon baselines; 200 posts hand-labelled; McNemar and bootstrap tests, chi-square tests across topics, Prophet forecasting. Developed from the cross-platform analytics idea in my Chevening application.
Aug 2017 - Jan 2022	BS Computer Science. National University of Computer and Emerging Sciences (FAST-NUCES), Pakistan. Courses included Probability and Statistics, Linear Algebra, Calculus I and II, Numerical Methods, Design and Analysis of Algorithms, Database Systems and Applied Machine Learning.

RESEARCH

Ikram, A. and Haig, E. "Is the Platform Part of the Measurement? A Protocol and Simulation Study for Cross-Platform Equivalence of Machine-Labelled Policy Sentiment." Paper, being posted to arXiv; ICWSM submission planned. Treats each platform as a survey mode and applies generalizability theory, measurement invariance and differential item functioning, and score linking to machine-labelled text. I built the simulation: with known ground truth, pooling platform scores at face value gives the wrong sign of daily aggregate sentiment on roughly a quarter of days, from measurement differences alone. Code: github.com/asadfix/platform-measurement, Zenodo DOI [10.5281/zenodo.23105129](https://doi.org/10.5281/zenodo.23105129).
Label validity at Artemis AI (2026). Found that two LLM labelling tiers agreed on 82.9% of sentiment labels but 0% of toxicity labels because of mismatched taxonomies, while the model reported high confidence; a September 2026 audit showed reported accuracies measured agreement with the labelling model rather than with people. Evaluation is being rebuilt around a human-labelled set. *No peer-reviewed publications to date.*

EXPERIENCE

2025 - present part-time	Co-Founder and Chief Technology Officer , Artemis AI, United Kingdom. Social-media intelligence for creators and businesses. - Lead a team of five across models, AI/ML, deployment, app development and design; take part in pitching. - Architecture agreed 6 Oct 2026 on AWS (London): Meta Graph API collectors; Apache Iceberg warehouse on S3 Tables and Athena; a three-tier labelling cascade (own NLP models, then DeepSeek, then Claude when confidence is low); image and video understanding (OCR, CLIP, Whisper); nightly scoring and live model services; 48 models in total. - Core collection and labelling run weekly today. In build for the beta on 15 December 2026: Art-E, an advisor on Amazon Bedrock; crisis alerts; 15 additional models. Hands-on on model accuracy; reading includes <i>Text as Data</i> (Grimmer, Roberts and Stewart) and LLM-annotation validation studies.
Dec 2024 - present part-time, remote	Data Engineer , M+C Saatchi Fluency, London. Social media analysis for global brands (consumer goods, automotive, sport, e-commerce, health) and UK government. - Social data collection, labelling and tagging, text and sentiment analysis with AI/ML, forum sourcing and pipeline architecture. - Designed and operate a fleet of 64 forum crawlers with a self-repairing LLM agent that diagnoses a broken source, regenerates the spider and tests it on live pages, secured with prompt fencing and approval gates; I approve each merge.

May 2023 - present contract, remote	- Operate a weekly pipeline of 2.2 million rows. Analyst on a twelve-month discourse study (100,845 Brandwatch and 170,542 Reddit mentions): BERTopic codebook, versioned LLM tagging. Data Engineer (contract) , Fix.com, Canada. Replacement-parts sites across the USA, Canada and Europe. - Built data architectures on Hostinger and open-source tools; OCR-based machine learning to parse parts catalogues from PDFs; models to backfill missing parts data. - Since June 2026 designed and run, alone, the company's crawler platform: about 10 million rows across 82 datasets by October 2026.
Jun - Jul 2025 contract, remote	Lead Data Engineer , VendueTech, Ireland. Led the Azure-to-AWS migration with an AWS partner and a team of three; designed the architecture for bidding prediction and data collection in the Irish market.
Sep - Dec 2024 contract, remote	Data Engineer , Prefe, Milan. Grocery e-commerce start-up. Migrated legacy collectors to Scrapy and helped design the machine-learning architecture.
Feb 2022 - Aug 2024 full-time	Associate Software Engineer, then Data Engineer, then Senior Data Engineer , Dubizzle Labs (Dubizzle, Bayut, Zameen, OLX), Lahore. Promoted twice in about two and a half years. - Owned a data-warehousing product end to end: data modelling, business rules and objectives. - Designed the company's first web-scraping architecture on AWS: crawlers for 15 markets, multilingual data in English, Arabic, Urdu and Hindi; led a team of three engineers in scraping and data warehousing. - Data modelling and machine learning for competitor-intelligence dashboards: what people discussed, how much volume each site carried, and what worked on each platform and in which region. - Helped junior colleagues into freelance and remote work; two got remote jobs.
Dec 2018 - Feb 2026 part-time, remote	Freelance data engineer , Fiverr. 300+ completed orders in web scraping, data engineering and analytics. Started during my BS; helped fellow university students start freelance careers.
Various dates advisory, remote	Web-scraping architecture advice and shorter engagements . A US travel-data start-up; businesses in the Netherlands, Lithuania and Germany; CXG, Dubai (freelance data engineering: production crawlers and a pytest QA framework).

SKILLS

Coding	Python (pandas, Scrapy, PyTorch, FastAPI, pytest), SQL, PySpark; Git, GitHub and CI. Daily work with AI coding assistants (Claude Code, OpenAI Codex): I write the specification and tests, review the changes and decide what merges.
Data wrangling and databases	Collection from APIs and the open web at scale; cleaning, deduplication, joins and validation; PostgreSQL, Redshift, Athena, DynamoDB, Apache Iceberg; data modelling and layered warehouses; AWS, Docker, Kubernetes.
Modelling and statistics	Hypothesis tests (McNemar, chi-square), bootstrap, Prophet forecasting; simulation with known ground truth; generalizability theory, measurement invariance and DIF, score linking; validating machine and LLM labels against human-labelled sets. Not yet trained in econometrics and causal inference or formal network science.
NLP and ML	Fine-tuning encoder classifiers (BERT, XLM-R, DistilBERT, DeBERTa-v3); LLM-assisted labelling (Claude, DeepSeek); OCR, CLIP and Whisper for images and video; BERTopic; LIME and integrated gradients.

AWARDS, COMMUNITY AND LANGUAGES

Chevening	Chevening Scholar 2024/25 (UK Government's global scholarship, which funded the MSc), awarded for helping students and junior colleagues into freelance and remote careers and for a proposed cross-platform social media analytics system to give government evidence of public opinion.
Public work	Public research code on GitHub and Zenodo (above). Invited by Zyte's developer-relations team to review a pre-release of its open-source agentic scraping toolkit (August 2026).
Languages	Urdu (native); English (working language; MSc taught and assessed in English, IELTS 6.5).

REFEREES

Dr Ella Haig, Associate Professor in AI, School of Computing, University of Portsmouth. Ella.Haig@port.ac.uk. MSc dissertation supervisor and co-author.

Dr Fahad Ahmad, Senior Lecturer, Centre for Cybercrime and Economic Crime, University of Portsmouth. Fahad.Ahmad@port.ac.uk. Taught and assessed me on the MSc.