

Asad Ikram

MSc Data Analytics (Distinction), University of Portsmouth · Chevening Scholar 2024/25
Applicant, Research Assistant / PhD Student, Geography of Code group, Corvinus University of Budapest
Lahore, Pakistan · +44 7482 376417 · asad.ikram53@gmail.com · linkedin.com/in/asad-ikram98 · GitHub: github.com/asadfix

PROFILE

Data engineer and researcher. Since December 2018 I have done remote software and data work for clients worldwide through Fiverr and on contracts for companies in Canada, Italy, Ireland, the United Arab Emirates, the United States and the United Kingdom. I write Python and SQL every day, build and repair large data collection and cleaning pipelines, and in 2026 direct AI coding agents that write most of the code I design, review and ship. My MSc research asks whether measurements built from online platform data mean the same thing across platforms; I want to study how AI and platforms change software work, using GitHub and other trace data.

EDUCATION

Sep 2024 - Sep 2025 awarded 8 Dec 2025	MSc Data Analytics, Distinction. University of Portsmouth, United Kingdom. Chevening Scholar. Credit-weighted average 70.5/100; master's project 78/100. Supervisor: Dr Ella Haig. Modules included Big Data and Machine Learning, Intelligent Data and Text Analytics, Explainable AI and Data Management (80). Taught and assessed in English. Dissertation: <i>Cross-Platform Sentiment Analysis of Public Reaction to UK Economic Policies.</i> About 279,000 public posts from X/Twitter, Reddit, YouTube and Quora, collected through official APIs and licensed feeds under university ethics approval (TETHIC-2025-111094). Fine-tuned BERT against BiLSTM and lexicon baselines (macro-F1 0.878, scored against lexicon-seeded labels); 200 posts hand-labelled by me, $\kappa \approx 0.78$ against VADER; McNemar and bootstrap tests, chi-square tests across topics, Prophet forecasting.
Aug 2017 - Jan 2022	BS Computer Science. National University of Computer and Emerging Sciences (FAST-NUCES), Pakistan. 130 credit hours, CGPA 2.75/4.00. Courses included Probability and Statistics, Linear Algebra, Calculus I and II, Numerical Methods, Design and Analysis of Algorithms, Database Systems and Applied Machine Learning.

RESEARCH

Ikram, A. and Haig, E. "Is the Platform Part of the Measurement? A Protocol and Simulation Study for Cross-Platform Equivalence of Machine-Labelled Policy Sentiment." **Manuscript, complete; not yet published.** Treats each platform as a survey mode and applies generalizability theory, measurement invariance and differential item functioning, and score linking to machine-labelled text. I built the simulation: in 500 replications with known ground truth, pooling platform scores at face value gives the wrong sign of daily aggregate sentiment on roughly a quarter of days, from measurement differences alone. Code: github.com/asadfix/platform-measurement, Zenodo DOI 10.5281/zenodo.23105129. *No peer-reviewed publications to date.*

EXPERIENCE

Dec 2024 - present part-time, remote	Data Engineer , M+C Saatchi Fluency, London. Consultancy analysing online conversation for brands and government. - Designed and operate 64 Scrapy crawlers collecting car-owner forum discussion for a global automotive brand on AWS, with a self-repairing LLM agent that diagnoses a broken source, regenerates the crawler and tests it on live pages behind a CI gate; I approve each merge. - Most of the fleet's new code, including a pipeline that builds new crawlers, is written by AI coding agents to my specifications and tests and reviewed by me before it merges. - Operate a weekly social-analytics pipeline for a global e-commerce client (2.2 million rows by 184 columns; Lambda, Fargate, PostgreSQL) with an 89-check data-quality sweep before each release. - Analyst on a study of twelve months of social discourse (100,845 Brandwatch and 170,542 Reddit mentions): BERTopic-derived codebook, LLM tagging with a prompt-version hash on every tag, three pilot runs on the same 500 rows compared before the full run.
May 2023 - present contract, part-time	Data Engineer (contract) , Fix.com, Canada (remote). Online replacement-parts retailer. Since June 2026 designed and built, alone, the company's crawler platform (about 10 million rows across 82 datasets by October 2026). Earlier: Scrapy pipelines over 200+ sources; Great Expectations validation that cut bad production data by 70%.

2025 - present part-time	Co-Founder and Chief Technology Officer, Artemis AI (formerly ViralData), United Kingdom (remote). Pre-launch social-media analytics start-up; I lead the technology in a team of eight; pilots from December 2026. - Designed an LLM labelling chain: Claude labels about 6,000 posts with an explanation of each; rules mined from the explanations become a prompt for DeepSeek, which labels about 12,000 more to train encoder classifiers (XLM-R, DistilBERT, DeBERTa-v3, Detoxify). The team first hand-labelled about 500 posts and comments to check both models. - Audited the labels: a September 2026 audit found that reported accuracies measured agreement with the labelling model, not with people; evaluation is being rebuilt around a human-labelled golden set. In progress: leading the build of Art-E, a retrieval-based advisor that answers each client's questions from that client's own data.
Sep 2024 - present part-time, remote	Shorter remote engagements. A US travel-data start-up, technical advisor, 2026: reviewed a travel-price data collection engine and measured its proxy network and traffic telemetry. CXG, Dubai, freelance data engineer, three engagements since Nov 2025: production crawlers and a pytest QA framework. VendueTech, Ireland, Lead Data Engineer, Jun - Jul 2025: led three engineers on real-time auction-data ingestion and an Azure-to-AWS migration (PySpark on EMR Serverless). Prefe, Milan, Data Engineer, Sep - Dec 2024: moved legacy collectors to Scrapy, cutting the failure rate by 60%; mentored two junior engineers.
Feb 2022 - Aug 2024 full-time	Associate Software Engineer, then Data Engineer, then Senior Data Engineer, Dubizzle Labs (Dubizzle, Bayut, Zameen, OLX), Lahore. Promoted twice in about two years. - Led a three-engineer team running 500+ scrapers across 15 markets on Kubernetes/EKS, collecting text in English, Arabic, Urdu and Hindi; redesigned the layered data warehouse, raising ETL throughput by 60%. - Built competitor-analysis dashboards from e-commerce and real-estate listings and online discussions that predicted revenue and showed what customers said.
Dec 2018 - Feb 2026 part-time, remote	Freelance data engineer, Fiverr. 300+ completed orders in web scraping, automation, ETL and data analysis; Level Two seller (first reached August 2021). Started during my BS and ran alongside later employment.

SKILLS

Coding	Python (pandas, Scrapy, PyTorch, FastAPI, pytest), SQL, PySpark; R (functional); Git, GitHub and CI. Daily work with AI coding agents (Claude Code, OpenAI Codex): I write the specification and tests, review the diff and decide what merges.
Data wrangling and databases	Collection from APIs and the open web at scale; cleaning, deduplication, joins and validation (Great Expectations, custom quality sweeps); PostgreSQL, Redshift, Athena, DynamoDB; layered warehouses; AWS (Glue, Step Functions, Fargate, Lambda, SageMaker), Docker, Kubernetes.
Modelling and statistics	Hypothesis tests (McNemar, chi-square), bootstrap, calibration, Prophet forecasting; simulation with known ground truth; generalizability theory, measurement invariance and DIF, score linking; validating machine and LLM labels against hand-labelled sets. Not yet trained in panel econometrics or formal network science.
NLP and ML	Fine-tuning encoder classifiers (BERT, XLM-R, DistilBERT, DeBERTa-v3); LLM-assisted labelling through the Claude and DeepSeek APIs; BERTopic; LIME and integrated gradients. I have not fine-tuned a generative model.

AWARDS, PUBLIC WORK AND LANGUAGES

Chevening	Chevening Scholar 2024/25, the UK Government's global scholarship, which funded the MSc.
Public work	Public code for the measurement manuscript on GitHub and Zenodo (above). Invited by Zyte's developer-relations team to review a pre-release of its open-source agentic scraping toolkit (August 2026).
Languages	Urdu (native); English (working language; MSc taught and assessed in English, IELTS 6.5).

REFEREES

Dr Ella Haig, Associate Professor in Artificial Intelligence, School of Computing, University of Portsmouth. Ella.Haig@port.ac.uk. MSc dissertation supervisor and co-author.

Dr Fahad Ahmad, Senior Lecturer, Centre for Cybercrime and Economic Crime, University of Portsmouth. Fahad.Ahmad@port.ac.uk. Taught and assessed me on the MSc.