CSH interview: easy guide

Tuesday 20 October, 8:40 pm Lahore time (4:40 pm UK) · Zoom, meeting ID 843 5332 6447 · about 10 minutes · recorded · in English

1. The cheat sheet

What they want to know: do you fit CSH, and can you do careful, independent research? Questions are about your experience and your interests. Nothing to study; just be clear and honest.

Your three messages

- **Your question:** when AI reads social media posts, do its labels mean the same thing on every platform? If not, how do we fix the comparison?
- **Your proof:** your first paper, with Dr Ella Haig, published on arXiv last week.
- **Your strength:** eight years working with social media data, so data is never your problem.

Your 60-second story (learn this)

> "My work has always been about social media data, and it grew step by step. At Dubizzle I first saw online text drive real business decisions, in four languages. That led to my Chevening idea: one system that compares platforms to show government what the public thinks. I built a small version in my MSc, and found my own check was circular. At Fluency I work with social media at scale for brands and UK government, and at my start-up, Artemis AI, I learned that AI labels can look fine exactly when they are wrong. My first paper turns these lessons into a method. A PhD at CSH is the next step: doing this properly, with network science."

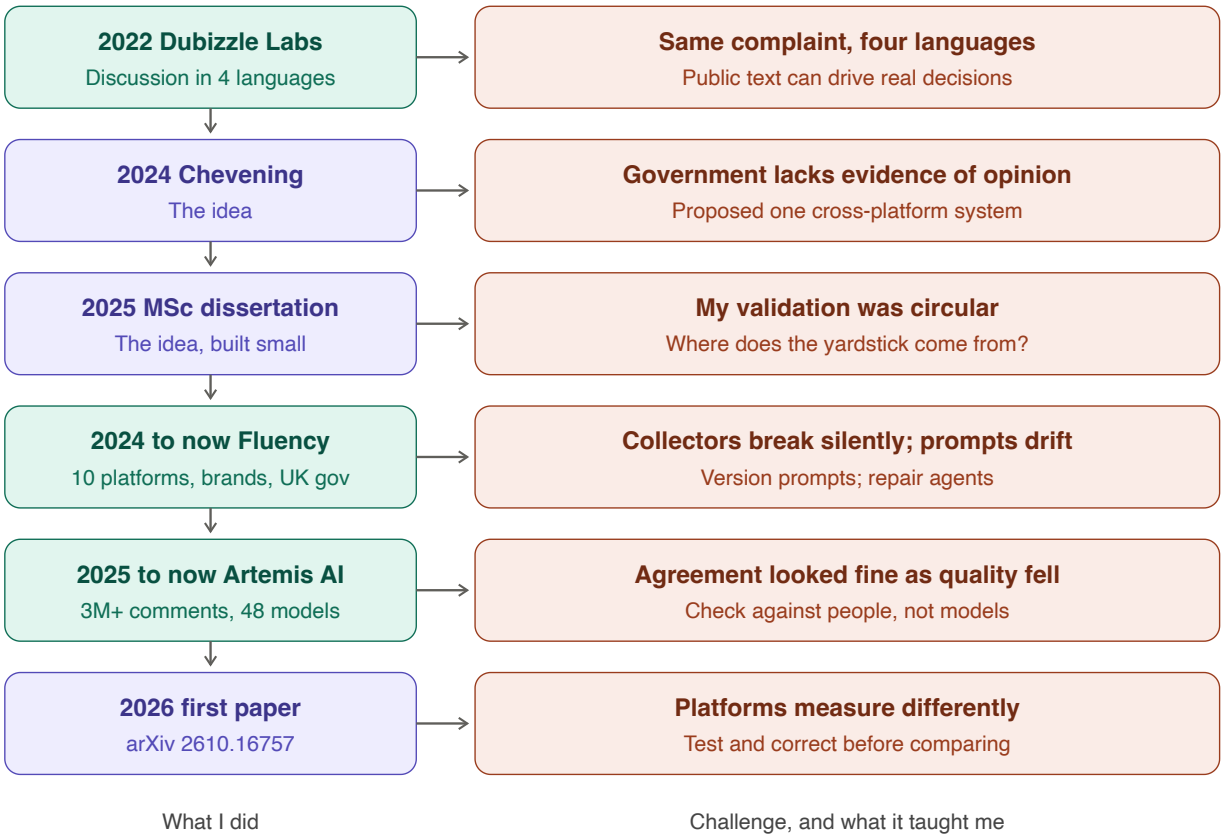
One line that ties it to the PhD

> "Every job showed me a new way social media measurements fail. The PhD is where I learn to test and fix them, and see how the fix changes what we conclude about narratives and polarisation."

Numbers to keep the same as your application

255,446 comments (MSc) **Team of eight at Artemis AI** **3 million+ comments** **arXiv 2610.16757** **MSc Distinction, Chevening 2024/25**

2. Your story, step by step



1. Dubizzle Labs (2022–2024)

Built the company's first web data system: 15 markets, four languages, a team of three. Dashboards showed what customers said online. **For the PhD:** you know how the same opinion reads differently across languages.

2. Chevening (2024)

Won for helping students into remote careers, and for an idea: one system comparing social media platforms to give government evidence of public opinion. **For the PhD:** your question is years old and useful to the public, not a trend.

3. MSc dissertation (2025)

Built the Chevening idea small: 255,446 comments from X, Reddit, YouTube and Quora about UK economic policy. Then you found the check on your AI model was circular. **For the PhD:** you can find the flaw in your own work and turn it into research.

4. M+C Saatchi Fluency (2024–now)

Social media from ten platforms every day, for big brands and UK government. You record which prompt produced every AI label, and built AI agents that fix broken data collectors. **For the PhD:** you can build clean, repeatable datasets from day one.

5. Artemis AI (2025–now)

Your start-up: 3 million+ comments, 48 models, a team of eight. The hard lessons:

- **Agreement can lie:** two AI models agreed more while both got worse. Always check against people.
- **Labels drift:** one measure fell from 77% agreement to 37% in weeks without anyone noticing.
- **“Normal” is hard:** out of 207,192 comments on busy posts, almost none were real storms. Each page needs its own baseline.
- **Words depend on the group:** banter and dialect get mistaken for abuse.

For the PhD: checking AI labels against people is exactly what CSH’s Alexander Hoyle works on.

6. Your first paper (October 2026)

All of this became one method: test whether platforms can be compared, and fix them if not. **For the PhD:** it is chapter one, already done.

3. Your paper, explained simply

Title: “Is the Platform Part of the Measurement? A Protocol and Simulation Study for Cross-Platform Equivalence of Machine-Labelled Policy Sentiment”. Ikram and Haig, arXiv 2610.16757, October 2026. Your first paper; a preprint, not yet peer reviewed.

The problem in one breath

Researchers and dashboards often collect posts from several platforms, let an AI score each post as positive, neutral or negative, and add everything up. But people write differently on each platform, so the AI may score the same feeling differently on each one. Then the total can be wrong.

The idea: treat a platform like a survey “mode”

Survey researchers know that people answer the same question differently by phone and on paper, and they test for this before comparing. Your paper says platforms are like modes too, so test them the same way.

What the paper does, in plain words

- **Step 1, same-event check:** compare platforms on the same news moment. You fixed 15 UK budget and economic events in advance (centred on the September 2022 mini-budget), before looking at any counts, so you couldn’t pick events that suit the answer.
- **Step 2, reliability check:** how much of the score comes from the platform, the AI model or the event, compared against posts labelled by people.
- **Step 3, “same meaning” test:** does “negative” behave the same on Reddit, X, YouTube and Quora? (Technical names: measurement invariance and DIF.)
- **Step 4, adjust for the crowd and timing:** different people post on each site at different times.
- **Step 5, one common scale:** convert every platform to one yardstick, like converting currencies before adding.
- **Step 6, a simulation:** a computer experiment where you know the true mood, to show what goes wrong if you skip steps 1 to 5.

The headline result

> Over 60 simulated days, simply adding platforms together showed the mood **pointing the wrong way on about 14 days** (nearly a quarter). Correcting each platform first: **under 1 day**. Repeated 500 times, so it is not luck.

Where it came from

The paper uses your dissertation as its real example: 255,446 comments on 6,650 posts about six UK economic topics (tax, cost of living, welfare, energy, housing, jobs). It openly reports that your model was checked against labels from the same word list it learned from, and that some dissertation numbers don’t match the records. Being open about this is a strength: it shows research honesty.

Its honest limits (know these)

- People choose their platforms, so differences mix “who posts” with “how the platform shapes posts”.
- Four platforms, UK only, English only.
- It measures expressed feeling, not support for a policy: a negative post about inflation isn’t opposition to a policy.
- It is a method plus a simulation; the full real-data study comes next, with 1,000–1,500 comments labelled by people and a second annotator.

If they ask about your paper

Question	Short answer
What did you personally do?	I built the simulation and the case study from my dissertation, and wrote it with my supervisor, Dr Ella Haig.
Is it published?	It is a preprint on arXiv; peer review is the next step.
What is new about it?	Others treated annotators or prompts as sources of error; nobody had treated the platform itself that way for machine-labelled text.
What next?	Run all six steps on fresh real data, and in a PhD extend it to AI models, time and how narratives spread.

4. Your dissertation, explained simply

“Cross-Platform Sentiment Analysis of Public Reaction to UK Economic Policies” · MSc Data Analytics, University of Portsmouth, September 2025 · supervisor Dr Ella Haig · Distinction (project mark 78/100) · ethics approved as low risk.

What it tried to do

Build the Chevening idea at small scale: collect what the UK public said online about economic policy on four platforms, measure the mood with AI, and see if past mood can forecast the next weeks.

Its four questions, in plain words

- Can modern AI read public mood about economic policy?
- How accurate are two AI models (BERT and BiLSTM)?
- Can past mood forecast short-term reactions?
- How does mood differ by topic and over time?

What you did

- **Data:** public comments from X/Twitter, Reddit, YouTube and Quora through official APIs and licensed archives. Use **255,446 comments**, the number the records confirm.
- **Labelling:** a word-list tool (VADER) gave first labels; then you trained BERT (a modern language model) and a BiLSTM (an older neural model) and compared them with simple baselines.
- **Checks:** statistical tests (McNemar, bootstrap), calibration (are the model’s confidence levels honest?), and LIME to show which words drove a decision.
- **Extras:** emotion analysis (anger and fear went with negative posts) and Prophet forecasting of daily mood.

What it reported

- BERT matched the labels 88% of the time, ahead of the BiLSTM (82%) and simple baselines; “neutral” was the hardest class.
- The mood index dipped sharply around the September 2022 mini-budget, and forecasts beat simple baselines.
- Platforms differed: YouTube looked more positive (explanatory videos), X had sharper negative spikes, Quora was sparse and noisy, and X dominated the data.

What you found later (the honest part)

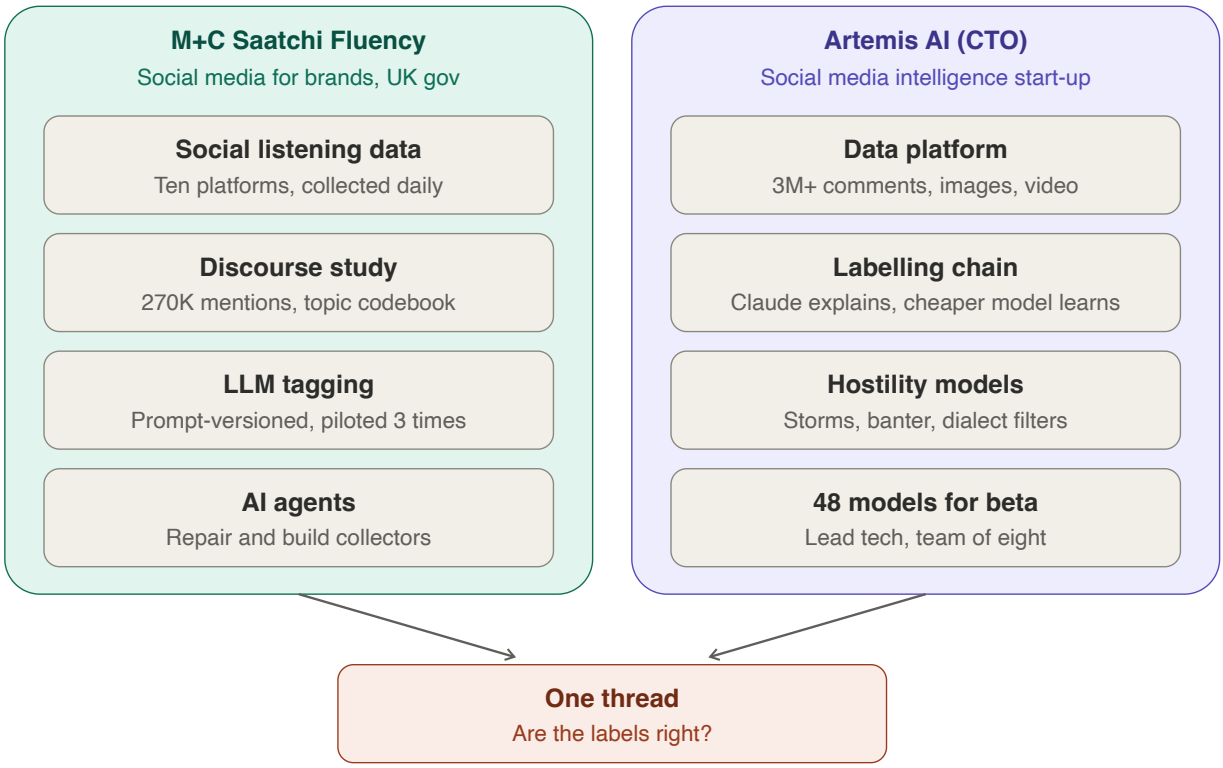
> “The 88% was agreement with labels that came from the same word list, so it proved less than it seemed. When I rechecked my records for the paper, a few numbers in the text also didn’t match. I report all of that openly in the paper, and it is exactly why the paper exists.”

If asked about topic results, don’t defend exact percentages; say the labels need checking against people first, which is what the follow-up study does.

5. All your experience, explained simply

Where and when	What you did, in plain words	Why it matters for this PhD
Fiverr 2018–2026	Started during your BS: 300+ web data projects for clients worldwide; Level 2 seller. Helped other students start freelancing.	Where you learned to collect web data and how public online data behaves.
Dubizzle Labs 2022–2024, Lahore	Built the company's first web data system: 15 markets, English, Arabic, Urdu, Hindi; led three engineers; dashboards on what customers said; promoted twice.	Multilingual data at scale; online text driving real decisions.
Fix.com 2023–now, Canada	Collection platform for a parts retailer (50M+ records), AI that reads PDF catalogues, data-quality checks.	Building clean, checked datasets.
Chevening 2024–25	UK Government scholarship, 70,000+ applicants; won for mentoring and the cross-platform public-opinion idea.	The origin of your research question.
MSc, Portsmouth 2024–25	Distinction; dissertation on four platforms (see the dissertation section).	Your first full research project.
Prefe, VendueTech, CXG 2024–26	Contracts in Italy, Ireland and Dubai: rebuilt data collectors, led three engineers through a cloud move, crawlers for luxury brands.	Leading people and delivering on deadlines.
M+C Saatchi Fluency 2024–now, London	Social media for Amazon, Ford, Nike, Reckitt and UK government: ten platforms daily, a twelve-month study of about 271,000 mentions, AI labels that record their prompt, AI agents that fix broken collectors.	Social media research at professional scale; reproducible AI labelling.
Artemis AI 2025–now	Your start-up, CTO, team of eight: 3M+ Facebook and Instagram comments, an AI labelling chain checked by people, hostility and storm detection, 48 models for the December beta.	Learned how AI labels fail in real life: the heart of your PhD question.
web-scraping-guide.com 2026	A free public guide to web data collection, including its legal and ethical limits.	You share knowledge and care about responsible collection.

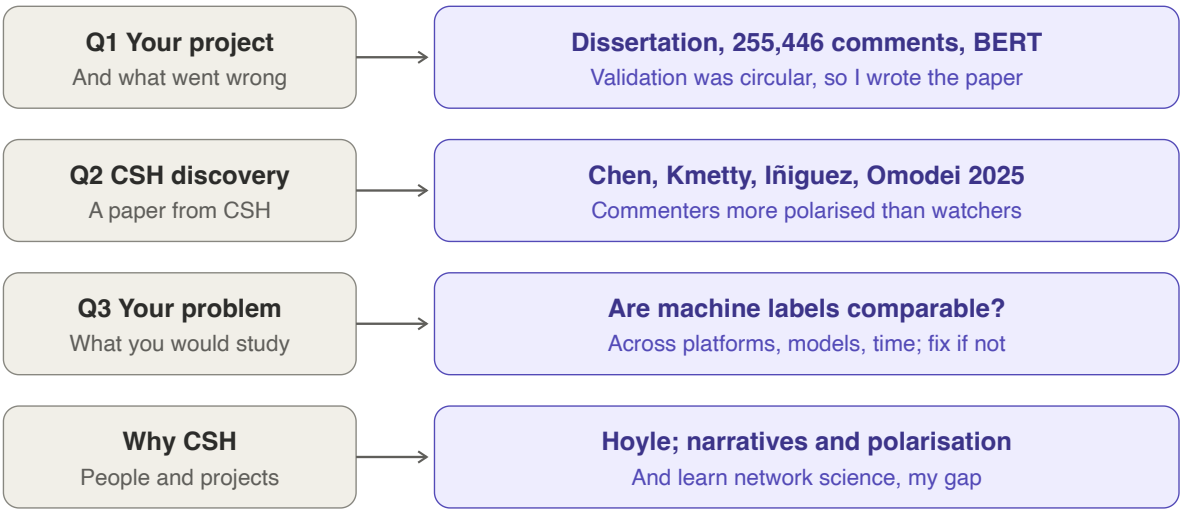
6. What you did at Fluency and Artemis AI



> “Brands and governments act on AI-labelled numbers that are rarely checked against people. My PhD would build the methods to check them.”

7. What you told CSH in your application

The interview recording is read together with these answers, so say the same things.



Answer 1: your project

Your MSc dissertation compared reactions to UK economic policy on four platforms. The AI model looked very accurate, but it had been trained and checked on labels from the same simple word list (VADER), so the check proved nothing. You also assumed “negative” means the same on every site. **Lesson:** always ask where your yardstick comes from.

Answer 2: a CSH paper you liked

Chen, Kmetty, Iñiguez and Omodei (2025): people who *comment* on political YouTube videos are more polarised than people who only *watch*. So studies built on comments make society look more divided than it is. **Your idea:** could a small group of people who donate their data be used to correct millions of public comments?

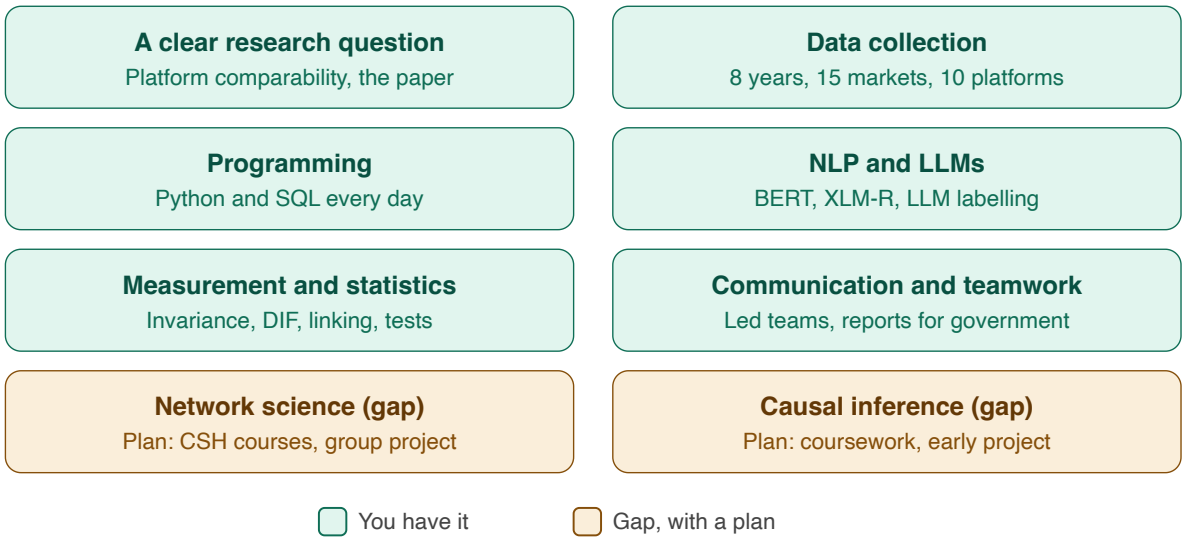
Answer 3: what you want to study

When can we compare AI-labelled opinions across platforms, AI models and time, and how do we fix them when we can’t? Four steps: extend your simulation; collect new posts around fixed events, labelled by several AIs and checked by people; test whether removing the “platform signal” removes real differences; then see how the fix changes how narratives seem to spread.

People and projects you named (all checked)

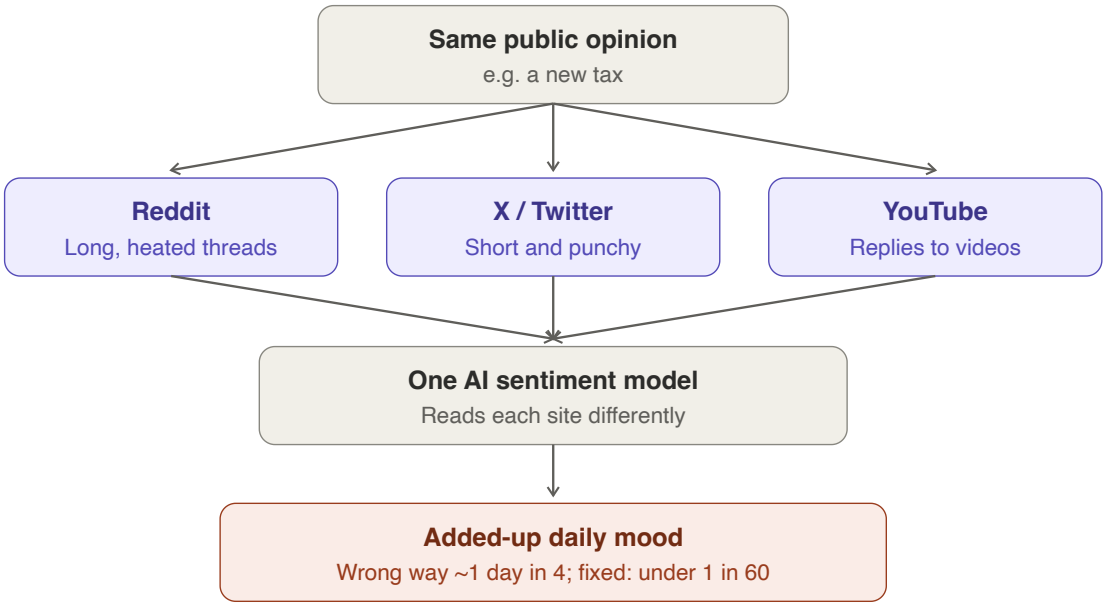
- **Alexander Hoyle**, CSH faculty: checking whether AI text measures really measure what they claim. Closest to you.
- **Elisa Omodei**, CSH faculty: co-author of the paper in Answer 2.
- Springboard projects: **Understanding narratives at scale**; **Polarization and minority representation on online platforms**.

8. What this PhD needs, and what you bring



Green: you already have it. Amber: you don't yet, and CSH teaches it, which is a good reason to do the PhD there.

9. Your paper in 20 seconds



> "People write differently on Reddit than on YouTube, so AI reads their opinions differently. Adding platforms together can flip the answer: in my simulation, the overall mood pointed the wrong way on nearly a quarter of days. My paper gives a checklist to test whether a comparison is fair, and a way to fix it."

10. Likely questions, short answers

Aim for 45–75 seconds. Say the main point first, give one example, then stop.

Question	Main points
Tell us about yourself	The 60-second story on page 1.
A research project, and what went wrong	STAR Story 1 (the circular check).
What have you built in industry?	STAR Story 4 (Fluency) or Story 5 (Artemis AI), then the one line in section 6.
Why a PhD, why now?	Every job showed a new way measurements fail; industry let you find problems but not study them; CSH gives network science and causal inference.
Why CSH?	Opinion online is a system of people, platforms and algorithms. Chen et al. and Hoyle's work match your question. You want to learn network science.
What would you study?	Answer 3 in three steps: simulation, new data checked by people, effect on how narratives spread.
Strengths and gaps	Strengths: data at scale, Python, AI labelling, careful checking. Gaps: network science, causal inference, with a plan.
Teamwork	STAR Story 3 (Dubizzle); add: lead the tech team of eight at Artemis AI.
A setback	STAR Story 1, focusing on the Action and Result.
Five years from now	Researching how online opinion is measured, so decisions rest on numbers checked against people.
How do you judge if data or a model is right?	STAR Story 5 or 6.
What drives you? / Why this topic?	STAR Story 2 (Chevening).
Questions for us?	"How are students matched with supervisors in the first year?"

11. STAR answer bank

STAR = **S**ituation (where you were), **T**ask (what had to be done), **A**ction (what *you* did), **R**esult (what happened, with a number). Say each in about 60 seconds, then add one line on what it means for the PhD.

Story 1: the circular check (research project, setback, honesty)

S	MSc dissertation at Portsmouth: public reaction to UK economic policy on X, Reddit, YouTube and Quora.
T	Build a reliable way to measure the public mood across platforms, and test it.
A	Collected 255,446 comments under ethics approval, trained BERT against baselines, ran statistical tests. Then I re-examined my own validation and found the model had been checked against labels from the same word list it learned from.
R	I reported the problem openly and turned it into my first paper with Dr Haig, which shows that ignoring platform differences gets the mood wrong on nearly a quarter of days.

For the PhD: I find flaws in my own work and turn them into research questions; that is the project I want to continue.

Story 2: Chevening (motivation, long-term commitment)

S	Working with online data in Pakistan, I saw that public bodies have little reliable evidence of what citizens think.
T	Win a fully funded UK master's to build that skill.
A	Proposed a cross-platform social media analytics system to give government evidence of public opinion, alongside my work helping students and junior colleagues into remote careers.
R	Selected as a Chevening Scholar from 70,000+ applicants; the proposal became my dissertation and then my paper.

For the PhD: the question is years old and useful to the public, not a new fashion.

Story 3: Dubizzle (leadership, multilingual data)

S	Dubizzle Labs had no system for collecting market and customer data across its regions.
T	Design the company's first one, covering 15 markets in English, Arabic, Urdu and Hindi.
A	Designed the architecture, led three engineers, added testing, and built dashboards on what customers said and how markets moved.
R	Dashboards used by decision-makers; production incidents down 45%; promoted twice in about two years.

For the PhD: I can lead a data effort and I know how the same view reads differently across languages, central to work on migration and minority discourse.

Story 4: Fluency (reproducible AI research at scale)

S	At M+C Saatchi Fluency, a twelve-month study of public discourse for a client: about 271,000 Brandwatch and Reddit mentions.

T	Code all of it reliably with AI, in a way that could be checked and repeated.
A	Worked on a topic codebook from BERTopic, LLM tagging where every label records the exact prompt version, and three pilot runs on the same 500 rows before the full run. Separately, built AI agents that repair broken data collectors, with a person approving each fix.
R	A traceable, repeatable dataset for the client, and collectors that no longer fail silently.

For the PhD: I can build reproducible corpora for "Understanding narratives at scale" from day one.

Story 5: Artemis AI, the agreement trap (judgement, validation)

S	At my start-up, we label millions of comments with AI models to train our own classifiers.
T	Check whether our labelling models were still reliable over time.
A	Compared the two models and noticed agreement had *risen* on three of five tasks while both were collapsing onto the most common answer. We moved our checks to fixed, human-checked labels and added monitoring for drift.
R	We caught a silent failure, and when one measure fell from 77% to 37% agreement, the monitoring made it visible. We are moving every model to checks against people before it reaches users.

For the PhD: validating AI labels against people is exactly Alexander Hoyle's area at CSH.

Story 6: Artemis AI, what is "normal"? (thinking like a researcher)

S	Clients wanted alerts for "comment storms": sudden waves of hostile comments.
T	Define a storm that is real, not a false alarm.
A	Had the team screen 207,192 comments on 369 unusually busy posts from February to July, and added filters so banter and dialect are not read as abuse.
R	Almost none were real storms aimed at the page: the hard part is each page's baseline, not the alarm itself.

For the PhD: polarisation and hostility need a baseline per community and platform, the same logic as commenters versus watchers in Chen et al.

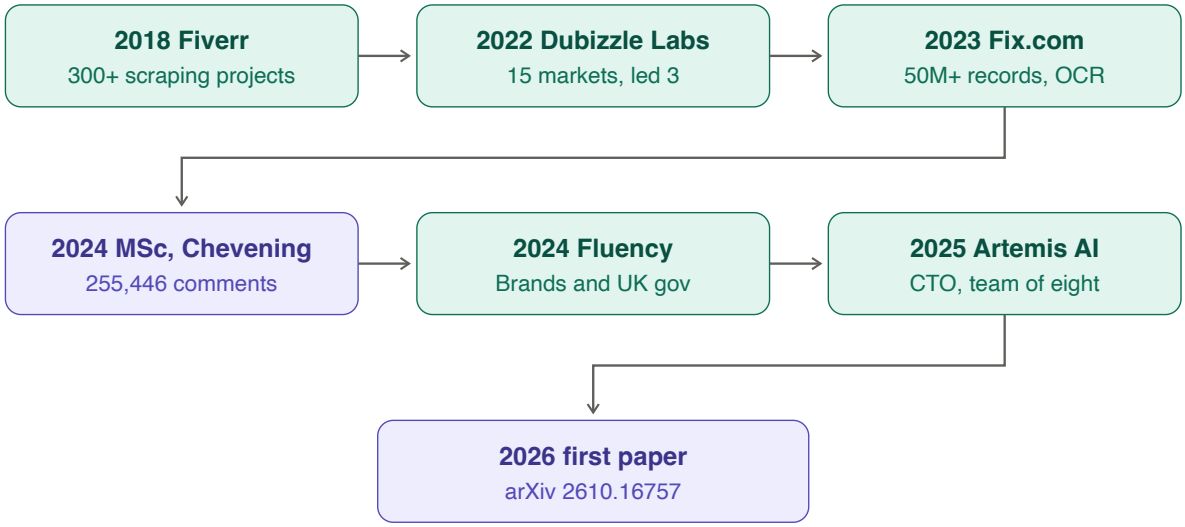
12. Glossary: hard words, made simple

Term	Meaning
Sentiment analysis	Using a computer to judge if a text is positive, neutral or negative.
VADER	A simple tool that scores sentiment by looking words up in a list.
BERT, BiLSTM	AI models that learn to read text; BERT is newer and usually better.
LLM (Claude, DeepSeek)	Large language models, here used to label posts instead of people.
Accuracy, F1	How often a model's labels match a reference. Only meaningful if the reference is trustworthy.
Circular validation	Checking a model against labels from the same source it learned from. It proves little.
Measurement invariance	Does a measure mean the same thing in different groups (here, platforms)?
DIF	"Differential item functioning": one label, like "negative", behaving differently in one group.
Generalizability study	Splits a score's variation into its sources: platform, model, event, people.
Score linking	Putting different scales onto one yardstick, like converting currencies.
Simulation with ground truth	A computer experiment where you set the true answer, so you can see how a method goes wrong.
Anchor events	Fixed news moments used to compare platforms on the same thing at the same time.
Data donation	People voluntarily share their own platform data (e.g. YouTube history) with researchers.
Polarisation	Groups moving apart; affective polarisation means disliking the other side.
Network diffusion	How ideas or narratives spread through connections between people or platforms.
Concept erasure	Removing one piece of information (such as which platform) from a model's view of a text.
Preprint / arXiv	A public research paper shared before peer review.

13. On the day, and what comes next

- Join at **8:35 pm**. Test camera and microphone. Quiet room, light in front of you, look at the camera.
- Keep this guide open but don't read from it. A short pause before answering is fine.
- Recording is normal; just go ahead.

If you pass: in 7–14 January you present a paper CSH chooses. Past choices were classic network papers: Watts and Strogatz 1998 (small worlds), Albert, Jeong and Barabási 2000 (network attacks), Clauset, Moore and Newman 2008 (missing links), Couzin et al. 2011 (group decisions). Start reading them after Tuesday. Finalists visit Vienna on 15–16 February; offers by mid-March. Referee letters must arrive by **1 November**.



Sources: CSH Graduate Program FAQ, DIS faculty list and springboard projects (csh.ac.at); CSH invitation email of 14 Oct; your CSH answers, CV and portfolio; Artemis AI shared notes and admin panel; PhD interview guides (AcademicPositions, Researcher.Life).