

Asad Ikram

Lahore, Pakistan · +44 7482 376417 · asad.ikram53@gmail.com · linkedin.com/in/asad-ikram98

Selection Committee

ELLIS Institute Finland

Doctoral student positions, autumn 2026 call

Re: Doctoral student positions at ELLIS Institute Finland, autumn 2026 call (requisition R47606)

Dear Members of the Selection Committee,

I run a distillation chain in production, and I would like to spend four years at ELLIS Institute Finland working out what the numbers it produces actually measure, with Eric Malmi as my first choice of supervisor. At ArtemisAI, which I founded and where I am Chief Technology Officer, Claude labels about 6,000 posts together with its reasoning, the decision rules in that reasoning are mined into a DeepSeek prompt, DeepSeek labels about 12,000 more, and six encoder classifiers learn from the result. It runs in production ahead of our December launch, and it taught me to distrust my own numbers. The figure our dashboard called accuracy was inter-model agreement with no criterion behind it: the two tiers agree on 82.9 percent of sentiment labels and 74.4 percent of emotion labels, while the toxicity model reports 96.7 percent average confidence. Versions v1 to v3 were scored against a human holdout and v4 onward against Claude's labels, so part of the apparent improvement is the reference moving, not the model. I was the person who could have caught that, and did not. Rebuilding the evaluation so that it cannot happen again is what turned a production bug into the question I want to work on.

When a large model labels the data and a cheaper one is distilled from it, the benchmark quietly becomes the teacher, and accuracy measures agreement, not validity. My interest at the Institute is how to validate LLM-based and distilled annotators used as measurement instruments for social data, across platforms, languages and model versions: what does accuracy mean when the reference is itself a model? The Institute fits because the question needs post-training and distillation on one side and computational social science on the other, and supervision across partner universities rather than one group.

My MSc in Data Analytics at the University of Portsmouth was awarded with Distinction in December 2025, credit-weighted average 70.5, project mark 78, as a Chevening Scholar selected from over 70,000 applicants. The dissertation analysed about 279,000 public posts on UK economic policy collected from X/Twitter, Reddit, YouTube and Quora, with a fine-tuned BERT classifier at macro-F1 0.878 and human validation at kappa around 0.78. The accuracy was not the finding. The finding was that the same policy event produced structurally different sentiment and emotion profiles on different platforms. With my supervisor Dr Ella Haig I turned that into a methods manuscript treating each platform as a survey mode, applying generalizability theory, invariance testing and score linking to machine labels. In a 500-replication simulation with known ground truth, a pooled index reports the wrong sign of the aggregate on roughly a quarter of days.

I would take that toolkit onto LLM annotation in four stages. First, the criterion: a human-anchored reference set stratified by platform and language, and a variance decomposition that separates model improvement from reference drift instead of summing them into one number. Second, equivalence: invariance testing and DIF ask whether a label means the same thing in two places before they are compared, and I would price what pooling costs when it does not. Third, drift: successive versions are alternate forms of one test, so score linking should keep a label series on one scale when the teacher changes, which is the failure I have had in production. Fourth, distillation itself: a student is trained to agree with its teacher, so agreement rises whether or not validity does, and the two belong in a report as separate numbers. I want the output to be a protocol a practitioner can run, because a method nobody else can run is not a method.

Eric Malmi is my first choice. His work at Google DeepMind on post-training, user signals and distillation is the mechanism my question is about, and his Aalto work on large language models in the social sciences and humanities is where it matters. Filip Ginter at Turku works on large web corpora and web register classification, which is where this question meets web scale: register is the likeliest axis along which a distilled annotator stops being comparable, and his corpora are where the second stage would find its hard cases. Corinna Coupette at Aalto aims to model, measure and manage complex systems, and her work on data-centric and responsible AI is where the first stage belongs. I would be glad to be matched with any of them, at the Institute or at a partner university.

Your criteria count industry work and software development alongside academic merit. Industry is the larger part of my record. Since December 2024 I have run the social listening pipeline at M+C Saatchi Fluency for UK Government departments and for Amazon, Ford and Nike: 68 collectors, 62 of them active, ten platforms daily, a governed 156-column master table, and 6.7 million items in the first six weeks. That record is why this question is not abstract to me. I own the instrument and I repair it on Mondays.

Two gaps, stated plainly. I have no publications: the manuscript with Dr Haig is complete, unpublished and not submitted, targeted at ICWSM 2027. My first degree is a BSc in Computer Science from FAST-NUCES, with a CGPA of 2.75 out of 4.00. Against them I would put the Distinction that followed and the production record built since. Two points of precision so that nothing here is read as more than it is: the models I fine-tune are encoder classifiers, and I use Claude and DeepSeek through APIs rather than fine-tuning generative models. I am available to relocate to Finland, my start date is flexible, and I do not speak Finnish.

I have spent seven years building the systems that turn what people say online into numbers, and the last two finding out how far one of those numbers travels before anyone asks what it was measured against. Four years is what it would take to answer that in a way others can use.

Yours faithfully,

Asad Ikram