

Asad Ikram

+44 7482 376417 · asad.ikram53@gmail.com · linkedin.com/in/asad-ikram98 · asad-ikram-mc.github.io/portfolio

Prof. Dr. Sonja Utz
Everyday Media Lab
Leibniz-Institut für Wissensmedien (IWM)
Schleichstraße 6, 72076 Tübingen

24 September 2026

Application: Researcher / PhD position, Everyday Media Lab, Ref. 1014-2026

Dear Professor Utz,

The AI system I built showed its average accuracy rising from 64.7 to 85.8 per cent over five versions, and I believed it. Your study with Anderl and colleagues (2024) found that people judged the same information more credible, and caught fewer of its errors, when a conversational agent presented it rather than static text. What misled me was a different cue, a precise number that kept rising, but the effect was the same, and I had every reason to know better. I am applying for the Researcher / PhD position in the Everyday Media Lab to study that belief in people who did not build the system.

My Chevening Scholarship (2024/25) rested on a simple argument: public decisions deserve better evidence about what people think. More of that evidence now comes from AI systems, so it is only as good as a reader's ability to tell when the system is wrong. I took it to the University of Portsmouth, where I was awarded an MSc in Data Analytics with Distinction in December 2025 (credit-weighted average 70.5, dissertation mark 78).

ArtemisAI, the social media analytics company I founded in 2025, is where I learned to distrust my own numbers. It is running in production ahead of our December launch and has processed over three million comments. Claude labels about 6,000 posts and DeepSeek about 12,000 more, both through their APIs, and six fine-tuned encoder classifiers learn from those labels. When I audited the rising accuracy figure, the first three versions had been scored against 97 human-labelled posts and the last two against Claude's labels, so part of the rise was agreement with another model rather than with people. On toxicity, the two labelling tiers used incompatible taxonomies and agreed on none of their labels, while the model reported 96.7 per cent average confidence. What stayed with me was my own judgement: I had built the system, and I still found its output credible on cues that told me nothing about whether it was correct.

The question is older than ArtemisAI. At Dubizzle Labs I led a team collecting text in English, Arabic, Urdu and Hindi across 15 countries, and saw that people write differently on forums and on social platforms. My MSc dissertation took that to data: an ethics-approved study of about 279,000 posts on UK economic policy from four platforms (X, Reddit, YouTube and Quora), with a fine-tuned BERT classifier (macro-F1 0.878) tested against BiLSTM and lexicon baselines, and a 200-post human-coded validation set (Cohen's kappa about 0.78). With my supervisor, Dr Ella Haig, I turned its central finding, that the same policy event looked structurally different on each platform, into a methods manuscript that treats each platform as a survey mode and adapts measurement invariance and score linking to machine-labelled sentiment. In a 500-replication simulation, a naively pooled index misstates a known event shift by roughly an eighth and gives the wrong sign on nearly a quarter of days. Writing it up, I found that my own 0.878 had been scored against lexicon-seeded labels rather than the human coders, and the manuscript says so.

Since December 2024 I have worked at M+C Saatchi Fluency, where I build data pipelines for clients including UK Government departments and, alongside the engineering, analyse what people say online for clients including Abbott, Reckitt, CFMOTO and a US women's soccer project. I can share how the work was done, not what it found. For the soccer project I built a Python pipeline over a year of US Twitter, forum and Reddit posts (over 100,000 consumer posts after filtering): account-level classification to separate fans from clubs, outlets and journalists, a theme codebook built from BERTopic topics that I read and merged by hand, and LLM tagging of every post, with a 500-row pilot checked by hand and repeated on the same rows to see where the model wavered. For Abbott I coded, in LLM-assisted stages, about 170,000 multilingual posts, forum threads and reviews from 12 platforms for emotions and need states, with a codebook that distinguishes health-professional, peer, family and brand sources.

For Reckitt I set product reviews against search queries on one shared needs codebook, comparing what people look for before buying with what they report afterwards. On a CFMOTO launch report, when the client questioned our method, we audited each brand's balance of news and consumer posts, set a minimum consumer base, re-ran the results on narrower subsets and compared two tagging runs, reporting as directional anything the checks did not support. Treating an LLM coder as an instrument that has to be tested before its output is believed is, I think, the right preparation for studying when people believe AI output, and it is how I would build stimuli of known accuracy.

Your paper with Goyanes and Gil de Zúñiga (2026) describes a cue route to trust in AI news through the AI heuristic, alongside an action route through seeing AI as a social entity. Observers also read a precise figure as a sign that the speaker is confident (Jerez-Fernandez, Angulo and Oppenheimer, 2014). People already overestimate how accurate LLM answers are (Steyvers et al., 2025), and a first-person hedge such as “I’m not sure” reduces reliance on them (Kim et al., 2024). My question is whether numeric precision acts as such a cue: whether a precise confidence figure raises credibility judgements independently of accuracy, and makes errors harder to detect, more than the same confidence stated in words.

Study 1 would be a preregistered online experiment extending the accurate versus partially inaccurate paradigm of Anderl et al. (2024), crossing how confidence is stated (verbally or as a precise figure, at matched levels, against a no-confidence control) with whether the content is accurate. The stimuli would be LLM summaries of real social media discussion, coded against human judgement so that every statement has a known accuracy, which lets the study use errors that real systems make rather than only planted ones. Study 2 would turn to incidental learning: participants meet the same summaries in a simulated professional feed, and an unannounced later test asks what they remember people said, including the wrong figures. Anderl, Gaiser and Utz (2026) found that work-related reading and posting predicted informational benefits at the same time point but not later, suggesting predominantly short-term benefits; Study 2 asks what persists when the information comes from an AI. I would reshape both designs with you.

My master's is in data analytics rather than psychology or communication science. My empirical work is quantitative but observational: I have not yet run an experiment with human participants, so I would draft and preregister Study 1 in my first months, within the lab's practice. I have no publications; the manuscript with Dr Haig is complete but not yet submitted. I do not speak German yet and would begin learning it on arrival.

I would accept the 75 per cent contract, relocate to Tübingen and agree a start date with you, and if my profile fits the agentic science position (1013-2026) better, I would be glad to be considered for it too. I have spent seven years building systems that turn what people say online into numbers, and the last two checking what those numbers were measured against. I would welcome the chance to discuss the fit with you. My CV, certificates and transcripts follow in this file.

Yours sincerely,

Asad Ikram

Chevening Scholar 2024/25 · MSc Data Analytics (Distinction), University of Portsmouth

Data Engineer, M+C Saatchi Fluency · Founder & CTO, ArtemisAI

asad.ikram53@gmail.com · +44 7482 376417

asad-ikram-mc.github.io/portfolio · linkedin.com/in/asad-ikram98

Asad Ikram

MSc Data Analytics (Distinction) · Chevening Scholar 2024/25 · Validated text analysis of social media data
Lahore, Pakistan · +44 7482 376417 · asad.ikram53@gmail.com · linkedin.com/in/asad-ikram98 · Portfolio: asad-ikram-mc.github.io/portfolio

PROFILE

I turn what people write on social media into numbers, and I test whether those numbers mean what they claim. For seven years I built the systems that collect and code public posts; for the last two I have been checking what they produce: validating automated coding against human coders, separating agreement between models from validity, and asking whether a score means the same thing on every platform. My Master's project was a designed, ethics-approved quantitative study of public text.

Research interests Credibility judgments of AI-generated information, especially when a precise confidence figure has no relation to correctness; incidental learning and informational benefits from social media and communicative AI at work; validity of automated text coding across platforms, models and time.

ANALYSIS METHODS

Content analysis	Automated content analysis validated against human coding; codebook design, piloting and revision; sentiment, emotion (NRC lexicon; a 24-emotion multi-label taxonomy), need-state, toxicity, topic and intent coding; topic modelling (BERTopic) to find candidate themes; LLM-assisted annotation checked by hand; multilingual text translated before coding. Classifiers are fine-tuned encoder models (BERT, XLM-R, DistilBERT, DeBERTa-v3).
Statistics	Inter-coder reliability (Cohen's κ); agreement versus criterion validity; McNemar tests, bootstrap confidence intervals, chi-square tests; probability calibration; simulation with known ground truth; checks of coverage, denominators and whether groups can be pooled; cross-market comparison and indexing; minimum-base and small-cell rules; sensitivity analyses; cross-tabulation of respondent-level survey data.
Measurement	Generalizability theory, measurement invariance, differential item functioning and score linking, specified in a protocol and tested by simulation; not yet fitted to collected data.
Research practice	Ethics approval and GDPR-compliant handling of public text; analysis anchors fixed in advance, before any counts were computed; documented, reproducible analysis pipelines.
Software	Python, PyTorch, SQL, PySpark; R (working knowledge). Data access through official APIs and licensed feeds (Meta Graph, Reddit, YouTube, Brandwatch, Sprinklr); AWS, Airflow, Docker.

EDUCATION

2024 - 2025 awarded Dec 2025	MSc Data Analytics, Distinction. University of Portsmouth, United Kingdom. Credit-weighted average 70.5/100 (Distinction threshold 70). Master's project 78/100, supervised by Dr Ella Haig. Modules included Intelligent Data and Text Analytics, Big Data and Machine Learning, and Explainable AI. 180 UK credits (conventionally 90 ECTS), taught and assessed in English.
2017 - 2021 awarded Jan 2022	BSc Computer Science. National University of Computer and Emerging Sciences (FAST-NUCES), Pakistan. Four years full time, 130 credit hours completed. CGPA 2.75 on a 0.00 to 4.00 scale.

RESEARCH

2025 Master's project	Cross-Platform Sentiment Analysis of Public Reaction to UK Economic Policies. • Study design. Ethics-approved quantitative study (Portsmouth TETHIC-2025-111094) of about 279,000 public posts and comments on UK economic policy from X/Twitter, Reddit, YouTube and Quora, collected through official APIs and licensed feeds under GDPR-compliant procedures. Human-annotated validation set of 200 posts balanced across platforms (reported inter-annotator agreement Cohen's $\kappa \approx 0.78$). • Classification and testing. Fine-tuned BERT classifier compared with BiLSTM by McNemar test ($p < 0.001$) and bootstrap intervals, alongside logistic regression and VADER baselines, with probability calibration. Chi-square tests of sentiment by policy topic; NRC emotion analysis linking anger and fear to negative sentiment in policy discussion. • Finding. The same policy event produced different sentiment and emotion profiles on different platforms, so a pooled index assumed, without testing, that the classifier measures the same construct everywhere. • Correction. Re-auditing for the manuscript, I found the headline macro-F1 of 0.878 was agreement with the lexicon-seeded training labels, not with the human set. The manuscript says so and specifies a larger stratified human criterion with a second annotator.
2026 manuscript	Ikram, A. and Haig, E. "Is the Platform Part of the Measurement? A Protocol and Simulation Study for Cross-Platform Equivalence of Machine-Labelled Policy Sentiment." Complete; not published and not submitted. • Treats each social platform as a survey mode and asks whether a classifier measures the same construct on each: a generalizability study against a human criterion, invariance and differential item functioning with platform as the grouping facet, and score linking. Fifteen anchor events fixed on 21 August 2026, before any window count was computed.

- 500-replication simulation with known ground truth: an index that pools platform scores at face value misstates a known event shift by roughly an eighth and reports the wrong sign of aggregate sentiment on nearly a quarter of days; the measurement-adjusted index gets the sign right on all but about one day in sixty.

No peer-reviewed publications to date.

EXPERIENCE

Feb 2025 - present	Founder and Chief Technology Officer, ArtemisAI Ltd (trading as ViralData), United Kingdom (remote). Social media analytics, running in production ahead of launch on 8 December 2026. • Text classification. The pipeline processes 111,000+ posts and 3.39 million comments from 30+ public pages. Six fine-tuned encoder classifiers code sentiment, emotion, toxicity, topic, language and intent. • Annotation design. Claude labels about 6,000 posts with its reasoning; rules mined from that reasoning prompt DeepSeek, which labels about 12,000 more; cases below 0.80 confidence return to Claude. • Evaluation design. Release gate specified as macro-F1 above 0.85 on both a 97-post human holdout and a 4,000-item stratified test set, followed by a human check of 200 random posts; weekly monitoring of shifts in output distributions. • Agreement is not validity. The two labelling tiers agree on 82.9% of sentiment, 74.4% of emotion and 0.0% of toxicity labels (a taxonomy mismatch), while the toxicity model reports 96.7% average confidence: a precise, confident number with no criterion behind it. • A moving reference. Reported average accuracy rose from 64.7 (v1) to 85.8 (v5), but v1 to v3 were scored against the human holdout and v4 onward against Claude's labels, so part of the gain is the reference moving rather than the model improving.
Dec 2024 - present	Data Engineer, consultant, M+C Saatchi Fluency, United Kingdom (remote) • Clients: UK Government departments, Abbott, Amazon, CFMOTO, Ford, Nike, Reckitt, a US soccer project. • Social discourse analysis. US women's soccer: a Python pipeline filtering 12 months of Twitter, forum and Reddit posts to over 100,000 consumer posts; accounts classified by who is speaking; BERTopic topics merged by hand into a codebook; every post LLM-coded for sentiment, themes and national-team versus club scope, checked by hand on a 500-post pilot; general public and fan communities reported apart. • Emotion and need-state analysis. Abbott: about 170,000 multilingual posts and reviews on infant formula from 12 platforms in Asia, Latin America and the Middle East, coded in LLM-assisted stages (translation, relevance, need states, 24 emotions, themes) and compared across markets, platforms and segments with small-cell suppression. Reckitt: over 33,000 US reviews LLM-coded on 16 functional needs, plus about 19,000 search keywords on the same framework, in an explorer tracing each figure to verbatims. • Method under challenge. CFMOTO 2027 launch report (eight brands): eight-theme sentiment codebook, composite metrics with written formulas; with the team, I answered client challenges with source-mix audits, a consumer-base minimum for all brands, re-tagging agreement checks, consumer-only and non-US sensitivity runs. Reckitt: a deep dive from respondent-level SPSS survey data, same filters and bases as a validated earlier deck; unsupported analyses dropped. • Infrastructure: 68 collectors (62 active) ingesting ten platforms daily, 6.7 million items in six weeks; one governed master table with documented filter rules and pre-delete backups.
Feb 2022 - Aug 2024	Senior Data Engineer, Dubizzle Labs (Dubizzle, Bayut, Zameen, OLX), Pakistan • Multilingual platform data in English, Arabic, Urdu and Hindi from 15 MENA countries. Led three engineers and 500+ collectors; redesigned the warehouse, ETL throughput up 60%. Promoted twice in two years.
2018 - present	Earlier and concurrent engagements • Fix.com, Canada, consultant, part time, May 2023 to present: 50M+ data points from 200+ sources; validation with schema-drift detection cut bad production data by 70%. • VendueTech, Ireland, Lead Data Engineer, Jun to Jul 2025: led three engineers, latency cut by 40%. Prefe, Italy, Sep to Dec 2024: failure rate cut by 60%, mentored two juniors. CXG, Dubai, Dec 2025 to Jan 2026. • Freelance data engineer, Dec 2018 to Feb 2026: 300+ five-star projects, 400+ websites, 50+ international clients.

AWARDS, PUBLIC WORK AND LANGUAGES

Chevening	Chevening Scholar 2024/25 (UK FCDO), selected from 70,000+ applicants in 160 countries; funded the MSc.
Public work	Author of web-scraping-guide.com, a free public reference on web data collection, its ethics and data quality.
Languages	Urdu (native). English: MSc taught and assessed in English; IELTS Academic 6.5 (March 2024). No German yet; I would begin learning on arrival.
Availability	Can start in November 2026 and relocate to Tübingen.

Only valid when viewed on port.ac.uk



UNIVERSITY OF PORTSMOUTH

Asad Ikram

has been awarded the degree of

**Master of Science in
Data Analytics**

with Distinction

having followed a programme of study approved by Academic Council

08 December 2025



Professor Graham Galbraith CBE,
Vice-Chancellor

Rebecca Di Pancrazio,
Academic Registrar



Student Number: 2291687 | Reference Number: 98611074-01-5ACA

Name of student **Asad Ikram**
Qualification **Master of Science in Data Analytics**
Award Classification **Distinction**
Date of Award **08 December 2025**
Mode of Study **Full Time**

University reference **2291687**
Date of birth **20 January 1998**
HESA reference **24100071550094239**

Awarding Institution **University of Portsmouth**
Teaching Institution **University of Portsmouth**

Record of Learning and Achievement

2024/2025	MSC DATA ANALYTICS - FHEQ Level 7	Mark	Credit	Decision
M32616	Masters Project	78	60	Pass
M33145	Programming for Data Analytics and AI	56	15	Pass
M33146	Big Data and Machine Learning Applications	67	30	Pass
M33147	Intelligent Data and Text Analytics	62	30	Pass
M33148	Data Visualisation and Exploratory Analytics	70	15	Pass
M33149	Data Management	80	15	Pass
M33516	Explainable AI	70	15	Pass
Credits gained for award			180	
Credits gained (ECTS)			90	

Information on the National Higher Education System can be obtained from National Academic Recognition Centre website:
www.enic-naric.net

For more information about this transcript please visit <https://myport.port.ac.uk/my-course/exams/after-your-assessments>

To check the validity of this transcript, scan the SmartVerify+ QR code below to confirm its authenticity or alternatively, you may contact the University of Portsmouth by the following methods:


By email: awards@port.ac.uk

By telephone: +44 (0) 23 9284 7745

Official Signatures:



Professor Graham Galbraith CBE,
Vice-Chancellor



Rebecca Di Pancrazio,
Academic Registrar

Not valid unless authorised by signature

Date Transcript Printed: 20 August 2026

END OF TRANSCRIPT



*Denotes Pass/Fail Module

Student Number: 2291687 | Reference Number: 61386216-01-UUAE

National University of Computer & Emerging Sciences



This is to certify that

Mr. Asad Ikram

has been admitted to the degree of

Bachelor of Science (Computer Science)

With all the honours, privileges, and responsibilities pertaining thereto.

Awarded in the city of Islamabad on the Thirty First day of January in the year 2022.

A stylized, handwritten signature in black ink, consisting of a large, looped 'A' followed by a horizontal line and a small flourish.

Rector



A handwritten signature in black ink, reading 'Wasim Sajjad' in a cursive style.

Chancellor



Student Name: **Asad Ikram**
Date of Birth: **January 20, 1998**
Univ. Reg. No: **17K-3943**

Roll No: **17K-3943**
Degree: **BS (Computer Science)**

Fall 2017

Code	CourseName	Crd	Pnt	Grd	Rmk
CL 101	Introduction to Computing Lab	1	2.00	C	
CS 101	Introduction to Computing	3	1.00	D	
EE 182	Basic Electronics	3	2.00	C	
MT 101	Calculus - I	3	3.00	B	
SL 101	English Language Lab	1	3.00	B	
SS 101	English Language	3	3.00	B	
SS 111	Islamic and Religious Studies	3	3.00	B	

Credits Attempted: 17 GPA: 2.41
Credits Earned: 17 CGPA: 2.41

Fall 2018

Code	CourseName	Crd	Pnt	Grd	Rmk
CL 103	Computer Programming Lab	1	1.33	D+	R-1
CS 201	Data Structures	3	0.00	F	
CS 211	Discrete Structures	3	1.33	D+	
EE 213	Comp. Organization and Assembly Lang.	3		W	
EL 213	Comp. Organization and Assembly Lang. Lat	1		W	
MG 103	Finance and Accounting	3		W	
MT 104	Linear Algebra	3	2.33	C+	

Credits Attempted: 43 GPA: 1.23
Credits Earned: 40 CGPA: 2.15

Fall 2019

Code	CourseName	Crd	Pnt	Grd	Rmk
CL 203	Database Systems Lab	1	4.00	A	
CL 205	Operating Systems Lab	1	2.33	C+	
CS 203	Database Systems	3	3.33	B+	
CS 205	Operating Systems	3	3.67	A-	
CS 301	Theory of Automata	3	3.33	B+	
MT 206	Probability and Statistics	4	3.33	B+	

Credits Attempted: 62 GPA: 3.38
Credits Earned: 62 CGPA: 2.69

Fall 2020

Code	CourseName	Crd	Pnt	Grd	Rmk
CS 302	Design and Analysis of Algorithms	3	2.00	C	
CS 303	Software Engineering	3	3.33	B+	
CS 406	Web Programming	3	3.33	B+	
CS 422	Human Computer Interaction	3	3.33	B+	
EE 204	Computer Architecture	3	2.33	C+	
SS 118	Psychology	3	3.67	A-	

Credits Attempted: 97 GPA: 3.00
Credits Earned: 97 CGPA: 2.76

Fall 2021

Code	CourseName	Crd	Pnt	Grd	Rmk
CS 4036	Software Testing	3	2.33	C+	
CS 4038	Data Mining	3	1.00	D	
CS 4040	Remote Sensing and GIS	3	3.33	B+	
CS 4092	Final Year Project - II	3	4.00	A	
MG 1001	Fundamentals of Management	3	3.00	B	
MG 1002	Marketing Management	3	2.67	B-	

Credits Attempted: 130 GPA: 2.72
Credits Earned: 130 CGPA: 2.75

Spring 2018

Code	CourseName	Crd	Pnt	Grd	Rmk
CL 103	Computer Programming Lab	1	0.00	F	
CS 103	Computer Programming	3	1.33	D+	
EE 227	Digital Logic Design	3	1.67	C-	
EL 227	Digital Logic Design - Lab	1	3.00	B	
MT 115	Calculus - II	3	2.67	B-	
SS 113	Pakistan Studies	3	3.33	B+	
SS 122	English Composition	3	3.00	B	

Credits Attempted: 34 GPA: 2.29
Credits Earned: 33 CGPA: 2.35

Spring 2019

Code	CourseName	Crd	Pnt	Grd	Rmk
CS 201	Data Structures	3	3.00	B	R-1
EE 213	Comp. Organization and Assembly Lang.	3	4.00	A	
EL 213	Comp. Organization and Assembly Lang. Lat	1	2.67	B-	
MT 207	Numerical Methods	3		W	

Credits Attempted: 47 GPA: 3.38
Credits Earned: 47 CGPA: 2.47

Spring 2020

Code	CourseName	Crd	Pnt	Grd	Rmk
CL 309	Object Oriented Analysis and Design Lab	1		S	
CL 323	Computer Networks Lab	1		S	
CS 309	Object Oriented Analysis and Design	3		S	
CS 323	Computer Networks	3		S	
CS 401	Artificial Intelligence	3		S	
CS 402	Compiler Construction	3		S	
MG 103	Finance and Accounting	3		S	

Credits Attempted: 79 GPA: N/A
Credits Earned: 79 CGPA: 2.69

Spring 2021

Code	CourseName	Crd	Pnt	Grd	Rmk
CS 4104	Applied Machine Learning	3	3.00	B	
CS 449	Professional Issues in IT	3	2.33	C+	
CS 491	Project - I	3	2.67	B-	
MT 207	Numerical Methods	3	2.67	B-	
SS 108	Technical and Business Writing	3	3.00	B	

Credits Attempted: 112 GPA: 2.73
Credits Earned: 112 CGPA: 2.75

CGPA Required: **2.00**
CGPA Earned: **2.75**

Credits Required: **130**

Degree Status: **Completed**

Credits Transferred: **0**
Credits Earned: **130**
Credits Completed: **130**

April 01, 2022

NATIONAL UNIVERSITY
of Computer & Emerging Sciences
Islamabad

Controller Examinations

Test Report Form

ACADEMIC

NOTE Admission to undergraduate and post graduate courses should be based on the ACADEMIC Reading and Writing Modules. GENERAL TRAINING Reading and Writing Modules are **not** designed to test the full range of language skills required for academic purposes. It is recommended that the candidate's language ability as indicated in this Test Report Form be re-assessed **after two years** from the date of the test. To find out more about IELTS, IELTS band scores and the CEFR levels, please visit ielts.org/scores

Centre Number	PK011	Date	28/MAR/2024	Candidate Number	507881
---------------	-------	------	-------------	------------------	--------

Candidate Details

Family Name	IKRAM
First Name(s)	ASAD
Candidate ID	3540423513455




Date of Birth	20/01/1998	Sex (M/F)	M	Scheme Code	Private Candidate
Country or Region of Origin					
Country of Nationality	PAKISTAN				
First Language	URDU				

Test Results

Listening	7.0	Reading	6.5	Writing	6.0	Speaking	7.0	Overall Band Score	6.5	CEFR Level	B2
-----------	-----	---------	-----	---------	-----	----------	-----	--------------------	-----	------------	----

Administrator Comments

Recognising organisations must verify this score at ielts.org/verify



Date	30/03/2024	Test Report Form Number	23PK507881IKRA011A
------	------------	-------------------------	--------------------