

Asad Ikram

+44 7482 376417 · asad.ikram53@gmail.com · linkedin.com/in/asad-ikram98 · asad-ikram-mc.github.io/portfolio

Prof. Dr. Sonja Utz
Everyday Media Lab
Leibniz-Institut für Wissensmedien (IWM)
Schleichstraße 6, 72076 Tübingen

24 September 2026

Application: Researcher / PhD position, Everyday Media Lab, Ref. 1014-2026

Dear Professor Utz,

The AI system I built showed its average accuracy rising from 64.7 to 85.8 per cent over five versions, and I believed it. Your study with Anderl and colleagues (2024) found that people judged the same information more credible, and caught fewer of its errors, when a conversational agent presented it rather than static text. What misled me was a different cue, a precise number that kept rising, but the effect was the same, and I had every reason to know better. I hold an MSc in Data Analytics (Distinction) from the University of Portsmouth and a BSc in Computer Science from FAST-NUCES, and I work as a Data Engineer at M+C Saatchi Fluency and as founder and CTO of ArtemisAI, a social media analytics company. I am applying for the Researcher / PhD position in the Everyday Media Lab to study that belief in people who did not build the system.

The problem first showed itself at Dubizzle Labs, where I was Senior Data Engineer from 2022 to 2024 and led three engineers collecting text in English, Arabic, Urdu and Hindi across 15 countries. People wrote differently on forums and on social platforms, so anything measured from their words depended on where it was collected. My Chevening Scholarship (2024/25) rested on a simple argument: public decisions deserve better evidence about what people think. More of that evidence now comes from AI systems, so it is only as good as a reader's ability to tell when the system is wrong. I took that argument to the University of Portsmouth (MSc awarded December 2025; credit-weighted average 70.5, dissertation mark 78).

My dissertation was an ethics-approved study of about 279,000 posts on UK economic policy from four platforms (X, Reddit, YouTube and Quora), with a fine-tuned BERT classifier (macro-F1 0.878) tested against BiLSTM and lexicon baselines, and a 200-post human-coded validation set (Cohen's kappa about 0.78). The same policy event looked structurally different on each platform. With my supervisor, Dr Ella Haig, I turned that finding into a methods manuscript that treats each platform as a survey mode and adapts measurement invariance and score linking to machine-labelled sentiment. In a 500-replication simulation, a naively pooled index misstates a known event shift by roughly an eighth and gives the wrong sign on nearly a quarter of days, yet it would reach a reader as one precise figure. Writing it up, I found that my own 0.878 had been scored against lexicon-seeded labels rather than the human coders, and the manuscript says so.

At Fluency, since December 2024, I have supported social listening pipelines for UK Government departments and Nike, built the ETL that cleans, translates and tags Amazon's social data, and run a self-repairing crawler fleet that gathers owner-forum discussion for Ford. On the analysis side, where I can share how the work was done but not what it found, I classified accounts for a US women's soccer project to separate fans from clubs, outlets and journalists, built a theme codebook from BERTopic topics I merged by hand, and LLM-tagged every post after a 500-row pilot checked by hand and re-run to see where the model wavered. For Abbott I coded, in LLM-assisted stages, about 170,000 multilingual posts and reviews from 12 platforms for emotions and need states; for Reckitt I set product reviews against search queries on one shared needs codebook. On a CFMOTO launch report, when the client questioned our method, we audited each brand's mix of news and consumer posts, set a minimum consumer base, re-ran the results on narrower subsets and compared two tagging runs, and reported as directional anything the checks did not support. Treating an LLM coder as an instrument that has to be tested before its output is believed is, I think, the right preparation for studying when people believe AI output.

ArtemisAI moved the problem from whether a figure is valid to why its readers believe it. It is running in production ahead of our December launch, and each feature turns a model output into advice that a small business owner acts on without being able to check it: flop detection, which warns when a page is at risk of a weak post; crisis detection, which raises alerts from toxicity scores when a comment storm starts; and a collaboration feature that applies network theory to a graph of page affinity to suggest whom to team up with. Underneath are NLP and OCR pipelines: six fine-tuned encoder classifiers (XLM-R, DistilBERT, DeBERTa-v3, Detoxify) code sentiment, emotion, toxicity, topic, language and intent across over three million comments, alongside OCR, Whisper transcription and a CLIP-based fusion classifier for images and video. The classifiers learn from labels that Claude (about 6,000 posts) and DeepSeek (about 12,000 more) produce through their APIs.

When I audited the rising accuracy figure, the first three versions had been scored against 97 human-labelled posts and the last two against Claude's labels, so part of the rise was agreement with another model rather than with people. On toxicity, the two labelling tiers used incompatible taxonomies and agreed on none of their labels, while the model reported 96.7 per cent average confidence. I had built the system, and I still found its output credible on cues that told me nothing about whether it was correct.

Your paper with Goyanes and Gil de Zúñiga (2026) describes a cue route to trust in AI news through the AI heuristic, alongside an action route through seeing AI as a social entity. Observers also read a precise figure as a sign that the speaker is confident (Jerez-Fernandez, Angulo and Oppenheimer, 2014). People already overestimate how accurate LLM answers are (Steyvers et al., 2025), and a first-person hedge such as "I'm not sure" reduces reliance on them (Kim et al., 2024). My question is whether numeric precision acts as such a cue: whether a precise confidence figure raises credibility judgements independently of accuracy, and makes errors harder to detect, more than the same confidence stated in words.

Study 1 would be a preregistered online experiment extending the accurate versus partially inaccurate paradigm of Anderl et al. (2024), crossing how confidence is stated (verbally or as a precise figure, at matched levels, against a no-confidence control) with whether the content is accurate. The stimuli would be LLM summaries of real social media discussion, coded against human judgement so that every statement has a known accuracy, which lets the study use errors that real systems make rather than only planted ones. Study 2 would turn to incidental learning: participants meet the same summaries in a simulated professional feed, and an unannounced later test asks what they remember people said, including the wrong figures. Anderl, Gaiser and Utz (2026) found that work-related reading and posting predicted informational benefits at the same time point but not later, suggesting predominantly short-term benefits; Study 2 asks what persists when the information comes from an AI. I would reshape both designs with you.

The answers would matter for any AI-generated summary of public opinion that professionals read at work, not only for my own product. In planning ArtemisAI's flop warnings we argued for showing owners bands rather than single scores, but I do not know whether that helps readers catch errors or only changes how sure they feel. That is an empirical question about credibility, and it belongs in a lab like yours rather than on a product roadmap. The PhD would give my work what it has lacked, tests with people rather than only against labels, and its answers would decide how systems like mine present their output. In return, I have built the dashboards whose precise figures I now want to test, and I know how they are computed and where they drift. I can generate stimuli from real LLM output, pretest them with the codebooks and human-coded checks I already write, build the simulated feed for Study 2 in Python, and adapt the known-ground-truth simulation from the manuscript to power analysis. I have spent seven years building systems that turn what people say online into numbers, and the last two checking what those numbers were measured against.

My master's is in data analytics rather than psychology or communication science. My empirical work is quantitative but observational: I have not yet run an experiment with human participants, so I would draft and preregister Study 1 in my first months, within the lab's practice, and learn the mixed-effects analysis such designs need. I have no publications; the manuscript with Dr Haig is complete but not yet submitted. I do not speak German yet and would begin learning it on arrival.

I would accept the 75 per cent contract and relocate to Tübingen, and if my profile fits the agentic science position (1013-2026) better, I would be glad to be considered for it too. My CV, certificates and transcripts follow in this file. Thank you for considering my application. I would welcome the chance to discuss the fit with you.

Yours sincerely,

Asad Ikram

Chevening Scholar 2024/25 · MSc Data Analytics (Distinction), University of Portsmouth

Data Engineer, M+C Saatchi Fluency · Founder & CTO, ArtemisAI

asad.ikram53@gmail.com · +44 7482 376417

asad-ikram-mc.github.io/portfolio · linkedin.com/in/asad-ikram98

Asad Ikram

MSc Data Analytics (Distinction) · Chevening Scholar 2024/25 · Validated text analysis of social media data
+44 7482 376417 · asad.ikram53@gmail.com · linkedin.com/in/asad-ikram98 · Portfolio: asad-ikram-mc.github.io/portfolio

PROFILE

At Dubizzle Labs, leading a team that collected text in four languages across 15 countries, I saw that people write differently on forums and on social platforms, yet their text was treated as comparable. As a Chevening Scholar I took that problem to the University of Portsmouth, into an MSc dissertation and a methods manuscript with Dr Ella Haig, then into client analysis at M+C Saatchi Fluency and into ArtemisAI, the company I founded, where I found my own system's accuracy figure rising partly because its reference had moved. I now want to study when people believe the numbers and outputs AI systems produce, which would also shape how ArtemisAI shows confidence to its users.

Fit with the Everyday Media Lab	Validating automated coding against human coders is how I would build stimuli of known accuracy and manipulation checks. McNemar tests, bootstrap intervals and simulation with known ground truth carry over to experimental data, and my Python pipelines could build a simulated professional feed for studying incidental learning.
Research interests	Credibility judgments of AI-generated information, especially when a precise confidence figure has no relation to correctness; incidental learning and informational benefits from social media and communicative AI at work; validity of automated text coding across platforms, models and time.

ANALYSIS METHODS

Content analysis	Automated content analysis validated against human coding; codebook design, piloting and revision; sentiment, emotion (NRC lexicon; a 24-emotion multi-label taxonomy), need-state, toxicity, topic and intent coding; topic modelling (BERTopic) to find candidate themes; LLM-assisted annotation checked by hand; multilingual text translated before coding. Fine-tuned encoder classifiers (BERT, XLM-R, DistilBERT, DeBERTa-v3).
Statistics	Inter-coder reliability (Cohen's κ); agreement versus criterion validity; McNemar tests, bootstrap confidence intervals, chi-square tests; probability calibration; simulation with known ground truth; checks of coverage, denominators and whether groups can be pooled; cross-market comparison and indexing; minimum-base and small-cell rules; sensitivity analyses.
Measurement	Generalizability theory, measurement invariance, differential item functioning and score linking, specified in a protocol and tested by simulation; not yet fitted to collected data.
Research practice	Ethics approval and GDPR-compliant handling of public text; analysis anchors fixed in advance, before any counts were computed; documented, reproducible analysis pipelines.
Software	Python, PyTorch, SQL, PySpark; R (working knowledge). Data access through official APIs and licensed feeds (Meta Graph, Reddit, YouTube, Brandwatch, Sprinklr); AWS, Airflow, Docker.

EDUCATION

2024 - 2025 awarded Dec 2025	MSc Data Analytics, Distinction. University of Portsmouth, United Kingdom. Credit-weighted average 70.5/100 (Distinction threshold 70). Master's project 78/100, supervised by Dr Ella Haig. Modules included Intelligent Data and Text Analytics, Big Data and Machine Learning, and Explainable AI. 180 UK credits (conventionally 90 ECTS), taught and assessed in English.
2017 - 2021 awarded Jan 2022	BSc Computer Science. National University of Computer and Emerging Sciences (FAST-NUCES), Pakistan. Four years full time, 130 credit hours completed. CGPA 2.75 on a 0.00 to 4.00 scale.

RESEARCH

2025 Master's project	Cross-Platform Sentiment Analysis of Public Reaction to UK Economic Policies. • Finding. The same policy event produced different sentiment and emotion profiles on different platforms, so a pooled index assumed, without testing, that the classifier measures the same construct everywhere. • Study design. Ethics-approved quantitative study (Portsmouth TETHIC-2025-111094) of about 279,000 public posts and comments on UK economic policy from X/Twitter, Reddit, YouTube and Quora, collected through official APIs and licensed feeds under GDPR-compliant procedures. Human-annotated validation set of 200 posts balanced across platforms (reported inter-annotator agreement Cohen's $\kappa \approx 0.78$). • Classification and testing. Fine-tuned BERT classifier compared with BiLSTM by McNemar test ($p < 0.001$) and bootstrap intervals, alongside logistic regression and VADER baselines, with probability calibration. Chi-square tests of sentiment by policy topic; NRC emotion analysis linking anger and fear to negative sentiment in policy discussion. • Correction. Re-auditing for the manuscript, I found the headline macro-F1 of 0.878 was agreement with the lexicon-seeded training labels, not with the human set. The manuscript says so and specifies a larger stratified human criterion with a second annotator.
2026 manuscript	Ikram, A. and Haig, E. "Is the Platform Part of the Measurement? A Protocol and Simulation Study for Cross-Platform Equivalence of Machine-Labelled Policy Sentiment." Complete; not published and not submitted.

- Treats each social platform as a survey mode and asks whether a classifier measures the same construct on each: a generalizability study against a human criterion, invariance and differential item functioning with platform as the grouping facet, and score linking.
 - 500-replication simulation with known ground truth: an index that pools platform scores at face value misstates a known event shift by roughly an eighth and reports the wrong sign of aggregate sentiment on nearly a quarter of days; the measurement-adjusted index gets the sign wrong on about one day in sixty.
- No peer-reviewed publications to date.

EXPERIENCE

Feb 2025 - present

Founder and Chief Technology Officer, ArtemisAI Ltd (trading as ViralData), United Kingdom (remote). Social media analytics for brand pages, running in production ahead of a December 2026 launch.

- **Agreement is not validity.** The two labelling tiers agree on 82.9% of sentiment, 74.4% of emotion and 0.0% of toxicity labels (a taxonomy mismatch), while the toxicity model reports 96.7% average confidence: a precise, confident number with no criterion behind it.
- **A moving reference.** Reported average accuracy rose from 64.7 (v1) to 85.8 (v5), but v1 to v3 were scored against the human holdout and v4 onward against Claude's labels, so part of the gain is the reference moving rather than the model improving.
- **Products.** Flop detection (warns when a page risks an underperforming post), crisis detection from toxicity scores, best posting times, post drafting, and a collaboration feature built on network theory (a networkx affinity graph of pages).
- **Models.** Six fine-tuned encoder classifiers on SageMaker code 111,000+ posts and 3.39 million comments from 30+ public pages (Meta Graph API): XLM-R (sentiment), DistilBERT (emotion), DeBERTa-v3 (topic, intent), Detoxify (toxicity) and a language classifier. They learn from a labelling chain run through APIs: Claude labels about 6,000 posts with its reasoning, rules mined from it prompt DeepSeek to label about 12,000 more, and cases below 0.80 confidence return to Claude.
- **Images and video.** A PyTorch fusion classifier over CLIP embeddings of image, OCR text and caption (macro-F1 0.93 on a 149-item gold set; production labels from Claude Haiku 4.5 vision); a hybrid OCR pipeline (token-F1 0.928); Whisper transcription.
- **Evaluation design.** Release gate of macro-F1 above 0.85 on both a 97-post human holdout and a 4,000-item stratified test set, then a human check of 200 random posts; weekly KL-divergence drift alarms.

Dec 2024 - present

Data Engineer, consultant, M+C Saatchi Fluency, United Kingdom (remote)

- **Method under challenge.** CFMOTO 2027 launch report (eight brands): eight-theme sentiment codebook and composite metrics with written formulas. With the team, I answered client challenges with source-mix audits, a consumer-base minimum for all brands, re-tagging agreement checks and consumer-only and non-US sensitivity runs; anything the checks did not support was reported as directional.
- **Social discourse analysis.** US women's soccer: over 100,000 consumer posts filtered in Python from a year of Twitter, forum and Reddit data; accounts classified by speaker; BERTopic topics merged by hand into a codebook; every post LLM-coded for sentiment and themes, hand-checked on a 500-post pilot.
- **Emotion and need-state analysis.** Abbott: about 170,000 multilingual posts and reviews on infant formula from 12 platforms in Asia, Latin America and the Middle East, coded in LLM-assisted stages (translation, relevance, need states, 24 emotions, themes) and compared across markets with small-cell suppression. Reckitt: over 33,000 US reviews and about 19,000 search keywords LLM-coded on one 16-need framework; a survey deep dive from respondent-level SPSS data, using the same filters and bases as a validated earlier deck.
- **Social listening pipelines** on AWS for UK Government departments, Amazon, Ford and Nike. UK Government: contributed to a communications-reporting pipeline (ingest, AI content tagging, comment processing, data-quality monitoring). Amazon: Sprinklr ETL (cleaning, deduplication, translation, tagging into Postgres and Tableau). Ford: 68 collectors on owner forums and other sources across ten platforms, with a self-healing repair agent.

Feb 2022 - Aug 2024

Senior Data Engineer, Dubizzle Labs (Dubizzle, Bayut, Zameen, OLX), Pakistan

- Led three engineers and 500+ collectors gathering text in English, Arabic, Urdu and Hindi from 15 MENA countries. This was where I first saw forum and social text treated as comparable. Promoted twice in two years.

2018 - present

Earlier and concurrent engagements

- **Fix.com**, Canada, part-time consultant since May 2023: 50M+ data points from 200+ sources.
- **VendueTech**, Ireland, Lead Data Engineer, Jun to Jul 2025: led three engineers. **Prete**, Italy, Sep to Dec 2024. **CXG**, Dubai, Dec 2025 to Jan 2026. **Freelance data engineer**, Dec 2018 to Feb 2026: 400+ websites, 50+ clients.

AWARDS, PUBLIC WORK AND LANGUAGES

Chevening

Chevening Scholar 2024/25 (UK FCDO), selected from 70,000+ applicants in 160 countries; funded the MSc.

Public work

Author of web-scraping-guide.com, a free public reference on web data collection, its ethics and data quality.

Languages

Urdu (native). English: MSc in English; IELTS Academic 6.5. No German yet; I would begin learning it on arrival.

Availability

Can start in November 2026 and relocate to Tübingen.

Only valid when viewed on port.ac.uk



UNIVERSITY OF PORTSMOUTH

Asad Ikram

has been awarded the degree of

**Master of Science in
Data Analytics**

with Distinction

having followed a programme of study approved by Academic Council

08 December 2025



Professor Graham Galbraith CBE,
Vice-Chancellor

Rebecca Di Pancrazio,
Academic Registrar



Student Number: 2291687 | Reference Number: 98611074-01-5ACA

Name of student **Asad Ikram**
Qualification **Master of Science in Data Analytics**
Award Classification **Distinction**
Date of Award **08 December 2025**
Mode of Study **Full Time**

University reference **2291687**
Date of birth **20 January 1998**
HESA reference **24100071550094239**

Awarding Institution **University of Portsmouth**
Teaching Institution **University of Portsmouth**

Record of Learning and Achievement

2024/2025	MSC DATA ANALYTICS - FHEQ Level 7	Mark	Credit	Decision
M32616	Masters Project	78	60	Pass
M33145	Programming for Data Analytics and AI	56	15	Pass
M33146	Big Data and Machine Learning Applications	67	30	Pass
M33147	Intelligent Data and Text Analytics	62	30	Pass
M33148	Data Visualisation and Exploratory Analytics	70	15	Pass
M33149	Data Management	80	15	Pass
M33516	Explainable AI	70	15	Pass
Credits gained for award			180	
Credits gained (ECTS)			90	

Information on the National Higher Education System can be obtained from National Academic Recognition Centre website:
www.enic-naric.net

For more information about this transcript please visit <https://myport.port.ac.uk/my-course/exams/after-your-assessments>

To check the validity of this transcript, scan the SmartVerify+ QR code below to confirm its authenticity or alternatively, you may contact the University of Portsmouth by the following methods:


By email: awards@port.ac.uk

By telephone: +44 (0) 23 9284 7745

Official Signatures:



Professor Graham Galbraith CBE,
Vice-Chancellor



Rebecca Di Pancrazio,
Academic Registrar

Not valid unless authorised by signature

Date Transcript Printed: 20 August 2026

END OF TRANSCRIPT



*Denotes Pass/Fail Module

Student Number: 2291687 | Reference Number: 61386216-01-UUAE

National University of Computer & Emerging Sciences



This is to certify that
Mr. Asad Ikram
has been admitted to the degree of
Bachelor of Science (Computer Science)

With all the honours, privileges, and responsibilities pertaining thereto.

Awarded in the city of Islamabad on the Thirty First day of January in the year 2022.

A stylized, handwritten signature in black ink, likely belonging to the Rector of the university.

Rector



A handwritten signature in black ink, likely belonging to the Chancellor of the university.

Chancellor



**NATIONAL
UNIVERSITY**
of Computer & Emerging Sciences
www.nu.edu.pk

Student Name: **Asad Ikram**
Date of Birth: **January 20, 1998**
Univ. Reg. No: **17K-3943**

Roll No: **17K-3943**
Degree: **BS (Computer Science)**

Fall 2017

Code	CourseName	Crd	Pnt	Grd	Rmk
CL 101	Introduction to Computing Lab	1	2.00	C	
CS 101	Introduction to Computing	3	1.00	D	
EE 182	Basic Electronics	3	2.00	C	
MT 101	Calculus - I	3	3.00	B	
SL 101	English Language Lab	1	3.00	B	
SS 101	English Language	3	3.00	B	
SS 111	Islamic and Religious Studies	3	3.00	B	

Credits Attempted: 17 GPA: 2.41
Credits Earned: 17 CGPA: 2.41

Fall 2018

Code	CourseName	Crd	Pnt	Grd	Rmk
CL 103	Computer Programming Lab	1	1.33	D+	R-1
CS 201	Data Structures	3	0.00	F	
CS 211	Discrete Structures	3	1.33	D+	
EE 213	Comp. Organization and Assembly Lang.	3		W	
EL 213	Comp. Organization and Assembly Lang. Lat	1		W	
MG 103	Finance and Accounting	3		W	
MT 104	Linear Algebra	3	2.33	C+	

Credits Attempted: 43 GPA: 1.23
Credits Earned: 40 CGPA: 2.15

Fall 2019

Code	CourseName	Crd	Pnt	Grd	Rmk
CL 203	Database Systems Lab	1	4.00	A	
CL 205	Operating Systems Lab	1	2.33	C+	
CS 203	Database Systems	3	3.33	B+	
CS 205	Operating Systems	3	3.67	A-	
CS 301	Theory of Automata	3	3.33	B+	
MT 206	Probability and Statistics	4	3.33	B+	

Credits Attempted: 62 GPA: 3.38
Credits Earned: 62 CGPA: 2.69

Fall 2020

Code	CourseName	Crd	Pnt	Grd	Rmk
CS 302	Design and Analysis of Algorithms	3	2.00	C	
CS 303	Software Engineering	3	3.33	B+	
CS 406	Web Programming	3	3.33	B+	
CS 422	Human Computer Interaction	3	3.33	B+	
EE 204	Computer Architecture	3	2.33	C+	
SS 118	Psychology	3	3.67	A-	

Credits Attempted: 97 GPA: 3.00
Credits Earned: 97 CGPA: 2.76

Fall 2021

Code	CourseName	Crd	Pnt	Grd	Rmk
CS 4036	Software Testing	3	2.33	C+	
CS 4038	Data Mining	3	1.00	D	
CS 4040	Remote Sensing and GIS	3	3.33	B+	
CS 4092	Final Year Project - II	3	4.00	A	
MG 1001	Fundamentals of Management	3	3.00	B	
MG 1002	Marketing Management	3	2.67	B-	

Credits Attempted: 130 GPA: 2.72
Credits Earned: 130 CGPA: 2.75

Spring 2018

Code	CourseName	Crd	Pnt	Grd	Rmk
CL 103	Computer Programming Lab	1	0.00	F	
CS 103	Computer Programming	3	1.33	D+	
EE 227	Digital Logic Design	3	1.67	C-	
EL 227	Digital Logic Design - Lab	1	3.00	B	
MT 115	Calculus - II	3	2.67	B-	
SS 113	Pakistan Studies	3	3.33	B+	
SS 122	English Composition	3	3.00	B	

Credits Attempted: 34 GPA: 2.29
Credits Earned: 33 CGPA: 2.35

Spring 2019

Code	CourseName	Crd	Pnt	Grd	Rmk
CS 201	Data Structures	3	3.00	B	R-1
EE 213	Comp. Organization and Assembly Lang.	3	4.00	A	
EL 213	Comp. Organization and Assembly Lang. Lat	1	2.67	B-	
MT 207	Numerical Methods	3		W	

Credits Attempted: 47 GPA: 3.38
Credits Earned: 47 CGPA: 2.47

Spring 2020

Code	CourseName	Crd	Pnt	Grd	Rmk
CL 309	Object Oriented Analysis and Design Lab	1		S	
CL 323	Computer Networks Lab	1		S	
CS 309	Object Oriented Analysis and Design	3		S	
CS 323	Computer Networks	3		S	
CS 401	Artificial Intelligence	3		S	
CS 402	Compiler Construction	3		S	
MG 103	Finance and Accounting	3		S	

Credits Attempted: 79 GPA: N/A
Credits Earned: 79 CGPA: 2.69

Spring 2021

Code	CourseName	Crd	Pnt	Grd	Rmk
CS 4104	Applied Machine Learning	3	3.00	B	
CS 449	Professional Issues in IT	3	2.33	C+	
CS 491	Project - I	3	2.67	B-	
MT 207	Numerical Methods	3	2.67	B-	
SS 108	Technical and Business Writing	3	3.00	B	

Credits Attempted: 112 GPA: 2.73
Credits Earned: 112 CGPA: 2.75

CGPA Required: 2.00
CGPA Earned: 2.75

Credits Required: 130

Degree Status: **Completed**

Credits Transferred: 0
Credits Earned: 130
Credits Completed: 130

April 01, 2022

NATIONAL UNIVERSITY
of Computer & Emerging Sciences
Islamabad

Controller Examinations

Test Report Form

ACADEMIC

NOTE Admission to undergraduate and post graduate courses should be based on the ACADEMIC Reading and Writing Modules. GENERAL TRAINING Reading and Writing Modules are **not** designed to test the full range of language skills required for academic purposes. It is recommended that the candidate's language ability as indicated in this Test Report Form be re-assessed **after two years** from the date of the test. To find out more about IELTS, IELTS band scores and the CEFR levels, please visit ielts.org/scores

Centre Number	PK011	Date	28/MAR/2024	Candidate Number	507881
---------------	-------	------	-------------	------------------	--------

Candidate Details

Family Name	IKRAM
First Name(s)	ASAD
Candidate ID	3540423513455




Date of Birth	20/01/1998	Sex (M/F)	M	Scheme Code	Private Candidate
Country or Region of Origin					
Country of Nationality	PAKISTAN				
First Language	URDU				

Test Results

Listening	7.0	Reading	6.5	Writing	6.0	Speaking	7.0	Overall Band Score	6.5	CEFR Level	B2
-----------	-----	---------	-----	---------	-----	----------	-----	--------------------	-----	------------	----

Administrator Comments

Recognising organisations must verify this score at ielts.org/verify



Date	30/03/2024	Test Report Form Number	23PK507881IKRA011A
------	------------	-------------------------	--------------------