

## Asad Ikram

Lahore, Pakistan · +44 7482 376417 · asad.ikram53@gmail.com · linkedin.com/in/asad-ikram98

Associate Professor Amir Hossein Payberah  
Department of Computing and Learning Systems  
School of Electrical Engineering and Computer Science  
KTH Royal Institute of Technology, Stockholm

### Re: Doctoral student in Feminist AI Addressing Gender-Based Violence (PA-2026-2872)

Dear Associate Professor Payberah,

In the labelling pipeline I run in production, two annotation tiers, Claude and DeepSeek, agreed on 0.0 percent of toxicity labels while the toxicity model reported 96.7 percent average confidence. The model was not failing. The two tiers were labelling against different toxicity taxonomies, and nothing in the confidence score could reveal it. That is the observation I would bring to this project: what an abuse detector finds depends on a definition someone chose, and the numbers we use to evaluate it are built to measure what it catches, not what it was never defined to see. Your project asks AI to make visible both explicit abuse and the structural silences around it. I would like to spend four years on the second half of that sentence, because it is where evaluation stops working.

**Why research.** I have spent seven years building the systems that turn what people say online into data, and the last two finding out that the most important questions about that data are not engineering questions. At ArtemisAI, which I founded and where I am Chief Technology Officer, Claude labels about 6,000 posts together with its reasoning, the decision rules in that reasoning are mined into a prompt for DeepSeek, DeepSeek labels about 12,000 more, and six fine-tuned encoder classifiers learn from the result, one of them a Detoxify-based toxicity model. It runs in production ahead of our December launch. Building it taught me that a reported accuracy is only as meaningful as the reference behind it: from version four onward our models were scored against Claude's labels rather than the human holdout, so part of the apparent improvement was the reference moving rather than the model. I want to do research because I can no longer build these instruments without asking what their outputs mean, and a doctorate is where that question gets the time it needs.

**Academic interests and previous studies.** My MSc in Data Analytics at the University of Portsmouth, awarded with Distinction in December 2025, was where the question became academic. My dissertation analysed about 279,000 public posts and comments on UK economic policy from X, Reddit, YouTube and Quora, with a fine-tuned BERT classifier at macro-F1 0.878 and human validation at Cohen's kappa around 0.78. The finding that mattered was not the accuracy: the same policy event produced structurally different sentiment and emotion profiles on different platforms, so one classifier was not measuring the same thing everywhere it read. With my supervisor Dr Ella Haig I developed that into a methods manuscript that treats each platform as a survey mode and applies measurement invariance and score linking to machine labels. In a simulation with known ground truth, a pooled index reports the wrong sign of the aggregate on roughly a quarter of days. My academic interest is the validity of machine labels for social data: whether a label means the same thing across platforms, languages, communities and model versions.

**How this relates to your project.** Gender-based violence online is where that question matters most. Whether a message counts as harassment depends on who wrote the definition, which community is speaking and in what language, and a detector trained on one definition will be confidently silent about harm that falls outside it. I would want to work on evaluation methods that surface those gaps: comparing annotations from different communities and from your feminist partner organisations to find where definitions diverge, testing whether a detector's errors fall disproportionately on particular groups or languages, and building datasets designed to capture what standard sampling leaves out. The project's plan to construct large-scale datasets fits the part of my work I am most confident in. At M+C Saatchi Fluency I maintain 68 collectors, 62 of them active, across ten platforms daily for UK Government departments, and at Dubizzle Labs I led a team running more than 500 collectors across 15 countries in English, Arabic, Urdu and Hindi. I am a native Urdu speaker, and a great deal of online abuse in South Asian languages is invisible to systems built for English, which is itself one of the silences this project describes.

**Gaps, stated plainly.** I have no publications yet: the manuscript with Dr Haig is complete but unpublished and not yet submitted, and I have attached it as a working paper together with my dissertation. I have no formal training in feminist theory, which this project is built on. I would expect to learn it from you and from the partner organisations rather than bring it, and I would rather say so than borrow its vocabulary. My work with graphs is limited to a networkx affinity graph, so knowledge graph construction would be new to me. And to be precise about my modelling: the models I fine-tune are encoder classifiers, and I use large language models through APIs rather than fine-tuning generative models myself.

**Future goals.** After the doctorate I want to keep working where machine learning meets accountability, building evaluation methods that let the organisations affected by automated systems, including civil society groups like the ones this project partners with, check what those systems can and cannot see. I am available to relocate to Stockholm, and my start date is flexible.

Yours sincerely,

**Asad Ikram**