Asad Ikram
Lahore, Pakistan · asad.ikram53@gmail.com · +44 7482 376417 · linkedin.com/in/asad-ikram98 · asad-ikram-mc.github.io/portfolio

Prof. Dr. Alexander Wuttke
Digitalization and Political Behavior, Geschwister-Scholl-Institut für Politikwissenschaft, LMU Munich

17 October 2026

Re: PhD Researcher in Computational Social Science, Emmy Noether Group "Predictably Paradoxical"

Dear Professor Wuttke,

Your migration interviews show that people at the same point on a scale can hold very different mental models, and that voice and chat draw different text out of the same people. Much of my last two years has gone into the second half of that problem in social media: when the medium changes how an opinion is written, can a model's labels still be compared? The position asks for Python and experience orchestrating LLM agents, text-as-data and NLP methods, and an interest in belief systems and democratic attitudes. I bring a methods paper on measuring opinion across platforms, LLM agents and labelling pipelines that I build and run in production, and eight years of collecting and analysing large multilingual text. I would like to use them to help your group map democratic belief systems that can be trusted across modes, languages and countries.

Your work. *AI Conversational Interviewing* (Wuttke, Lang, Klamm, Würschinger and Kreuter, 2026) shows what a battery from the German Longitudinal Election Study cannot: not only where partisans stand on migration, but what comes to mind first, with FDP supporters opening on the labour market and Die Linke supporters on humanitarian concerns. The paper is also candid about two things I recognise. Voice interviews produce longer, more conversational answers and text interviews shorter, denser ones, with chat replies shrinking after about the tenth turn, so the mode comparisons describe systems as deployed rather than a pure modality effect. And about a quarter of survey and interview records could not be linked because respondents typed their IDs by hand. The first is a measurement problem and the second a pipeline problem, and I have worked on both.

Python skills and experience working with LLMs. At M+C Saatchi Fluency in London, where I build social media data systems for brands and UK Government communications teams, I designed a self-healing agent: Claude works through a set of tools on a server I built, diagnoses a broken forum crawler, rewrites it and tests the fix on the live site before anything ships. A second agent builds new crawlers. Both run to an agenda and explicit ground rules, with prompt fencing and approval gates, and I decide what merges, which is the same discipline an AI interviewer needs to stay neutral and on protocol. As co-founder and CTO of Artemis AI, a social media intelligence start-up, I lead a team of five; our labelling escalates from our own models to DeepSeek and then Claude when confidence is low. I work in Python every day.

Text-as-data methods and natural language processing. My MSc dissertation in Data Analytics at the University of Portsmouth (Distinction, Chevening Scholar), supervised by Dr Ella Haig, analysed public reaction to UK economic policy in more than 250,000 posts from X, Reddit, YouTube and Quora, comparing a fine-tuned BERT model with BiLSTM and lexicon baselines against 200 hand-labelled posts. At Fluency I was an analyst on a twelve-month discourse study of 100,845 Brandwatch and 170,542 Reddit mentions, using a BERTopic codebook and versioned LLM tagging that people then checked: the same LLM-first, human-validated procedure your team used to code topics and frames. At Artemis AI I fine-tune encoder classifiers (XLM-R, DistilBERT, DeBERTa) for toxicity, intent and topic, tested against human-labelled sets.

Interest in belief systems and democratic attitudes. My research began with a political question: what does the public think of the policies made in its name, and can online text tell us? The dissertation became a paper with Dr Haig, *Is the Platform Part of the Measurement?* (arXiv, 2026), which treats each platform as a survey mode and applies generalizability theory, measurement invariance and score linking to machine-labelled opinion. In a simulation I built, pooling platform scores at face value gives the wrong sign of aggregate sentiment on nearly a quarter of days. My Chevening application proposed a cross-platform system to give government better evidence of public opinion. What draws me to your group is the step beyond sentiment: not whether people approve of something, but how their reasons hang together, and which of them hold up when they are asked a follow-up.

A question I could contribute to. Before belief systems are compared across modes and across your thirteen countries, are the topic and frame labels behind them invariant? If voice produces looser, longer reasoning and chat shorter, denser claims, a frame such as "causal reasoning" may be coded more often in one mode for the same underlying view, and a democratic attitude may look well-founded or superficial partly because of the interviewer

system. I would treat mode, language and country as facets, test the coding scheme for invariance and differential item functioning against human-coded subsamples, and only then compare. On the engineering side I would make linkage and turn-level logging automatic, so that no interview is lost to a typed ID. I have built text pipelines in English, Arabic, Urdu and Hindi, and I know how differently the same stance can read across languages.

Why a PhD, and why now. Industry taught me to build systems that collect and label opinion at scale; it did not teach me to show what those data say about people. I want that training in a group that takes measurement as seriously as I have learned to, and I would carry the results back into how AI tools for public opinion are built.

What I do not yet bring. I do not have a degree in political science or psychology, and my knowledge of political psychology comes from reading, not training. I have not designed a survey or run an experiment with human participants. My agents work through tools rather than talk to people, so I have not yet built a multi-turn conversational interviewer, and I have no peer-reviewed publication. I would close these gaps through the institute's courses and by working on the group's interview studies from the first month. I do not yet speak German and would start learning it on arrival.

Practicalities. I can start on 1 January 2027 and would relocate to Munich to work on site. As a writing sample I attach the arXiv paper, for which I built the simulation and the case study from my dissertation. My referees are Dr Ella Haig, Associate Professor in AI, University of Portsmouth, my supervisor and co-author (Ella.Haig@port.ac.uk), and Dr Fahad Ahmad, Senior Lecturer, University of Portsmouth, who taught and assessed me on the MSc (Fahad.Ahmad@port.ac.uk). Thank you for considering my application.

Yours sincerely,
Asad Ikram