Asad Ikram

Lahore, Pakistan · +44 7482 376417 · asad.ikram53@gmail.com · linkedin.com/in/asad-ikram98

Associate Professor Kim Andersen
Centre for Journalism, Department of Political Science and Public Management
University of Southern Denmark, Odense

Re: PhD student with a strong computational profile, Navigating Negativity: What is a Healthy News Diet? (SP1)

Dear Associate Professor Andersen,

I am applying for the PhD position in SP1. I hold an MSc in Data Analytics with Distinction from the University of Portsmouth, awarded in December 2025, and a BSc in Computer Science, and I have spent seven years building and maintaining large-scale social media data collection systems in industry. I am available to relocate to Odense and to start on 1 November 2026. On appointment I would step down from my role as CTO of ArtemisAI and end my consultancy at M+C Saatchi Fluency, so the post is full time and there is no competing claim on my hours and no commercial interest that could touch this project's data.

What draws me to SP1 is a puzzle in the negativity literature that I think measurement is currently hiding. The negativity bias in news is one of the better-established findings in political communication, and yet your preregistered work with Shehata finds a preference for positive news, and Denmark now reports one of the sharpest rises in news avoidance in Europe while remaining a high-trust, high-consumption media system. Those three things do not sit comfortably together. One possible reconciliation is behavioural: people say and choose differently, and avoidance is intentional in some cases and incidental in others, which is the distinction Skovsgaard and you drew. But another possibility is that the quantities being compared are not the same quantity. Negativity is not one thing, and Lengauer, Esser and Berganza separate it into axes that a single score collapses. If an outlet's negativity, a platform's negativity and a person's feed negativity are measured by one instrument that reads those three text types differently, then part of what looks like a difference in the world is a difference in the ruler. What I would want to know at the end of three years is concrete: how much more negative, if at all, is a Danish person's actual feed than the outlets that feed it, and is the answer the same for the tabloid reader and the broadsheet reader. Nobody can currently say, because the measurement needed to say it has not been built.

That is a measurement problem, and it is the one my own research is about. My MSc dissertation, supervised by Dr Ella Haig, analysed roughly 279,000 posts on UK economic policy across Twitter, Reddit, YouTube and Quora, with a fine-tuned classifier at macro-F1 0.878 and human validation at Cohen's kappa around 0.78, collected under Portsmouth ethics approval TETHIC-2025-111094 and GDPR-compliant procedures. The accuracy was not the finding. The finding was that the same policy event produced structurally different sentiment and emotion profiles on different platforms. Dr Haig and I developed that into a methods manuscript that treats each platform as a survey mode and adapts the psychometric equivalence toolkit, generalizability theory, measurement invariance and differential item functioning, and score linking, to machine-labelled text. In a simulation with known ground truth, an index pooling platform scores at face value reports the wrong sign of the aggregate on roughly a quarter of days. I want to be careful about what is new there: that automated measures lose validity out of domain is well established in this field, and what we add is the survey-mode framing and the linking machinery, with a quantified cost of assuming equivalence. The manuscript is complete but unpublished and not submitted, it is on English-language material, and it has not been reproduced on Danish data. I attach it with a statement of my contribution so it can be assessed as documented work.

The second reason I am applying is that the hub is a collection problem I have solved versions of for seven years. At M+C Saatchi Fluency I run sixty-eight collectors ingesting ten platforms daily into a governed 156-column warehouse, with an automated repair layer that diagnoses a broken source, regenerates the extraction logic and tests it against live pages before shipping. Two failure modes decide whether SP1 gets data at all. Download packages change structure without notice, and their folder and key names change with the account's language setting, which is decisive in Denmark where donors' phones will split between Danish and English; the standing defence is golden fixtures per platform and locale, versioned parsers stamped onto every record, and canary donations from project-owned accounts. And donated entries are identifiers and links rather than text, so a negativity classifier cannot read them until every item is resolved back to its content at scale, which Hase and colleagues name as the blocker that makes content analysis of donated feeds unfeasible. Both are engineering problems with measurement consequences: resolution fails for deleted and restricted items, which biases the recovered feed toward what is still live, and that is a survivorship bias correlated with the outcome.

ArtemisAI is where I learned to distrust my own numbers. I founded it and designed its architecture: six fine-tuned encoder classifiers for sentiment, emotion, toxicity, topic, language and intent, trained on a labelling pipeline where a strong model labels a seed set with its reasoning, decision rules are mined from that reasoning for a cheaper model, and low-confidence cases route back for review. On top sit the products clients wanted: predicting underperforming posts, posting windows, drafting, crisis alerts, and an affinity graph where I applied network theory from the literature to see whether it held in production. Then I read our own code and found that the dashboard's accuracy figure was inter-model agreement with no criterion behind it: our two annotation tiers agree on 82.9 per cent of sentiment labels and 74.4 per cent of emotion labels, while the toxicity model reports 96.7 per cent average confidence. I was the person who could have caught that and did not, for two years, because it never occurred to me to ask what the number was against. I have since put a human reference set behind the reporting. But it taught me the thing I would bring here: a tag on a post becomes a share-of-voice number, becomes a slide, becomes a communications strategy, and by the time it is a strategy nobody remembers the first step was a model with an error rate. In this project that chain ends in public advice about how much negative news is healthy, which is a considerably better reason to get the first step right.

On qualifications, my master's is a UK one-year taught MSc of 180 UK credits, conventionally mapped to 90 ECTS, with Distinction, a dissertation mark of 78 and a credit-weighted average of 70.5. It follows a four-year, 130-credit-hour BSc in Computer Science, so the combined study load exceeds a three-plus-two structure. I have written to the PhD coordinator for a view on equivalence and will forward the reply; both transcripts are attached so the assessment can rest on documents.

On Danish, I do not speak it and I am not proposing to hand the construct to other people. The codebook is mine to develop with the team, the reliability targets and the adjudication rules are mine to set, and I would work with Danish material from month one through a bilingual protocol: native-speaker coders produce the gold standard, and I sit inside the disagreements with them, because the cases two trained coders split on are where a codebook is actually written. Error analysis, which is most of the instrument work, does not require fluency; it requires being able to see which cases the model loses and asking a native speaker why. I would begin Danish tuition on arrival, and I would expect to be able to read the material unaided well before the coding for WP3 begins.

On teaching, I have none in a university setting. I led three-engineer teams at Dubizzle Labs and VendueTech, mentored two junior engineers at Prefe, and I write and maintain web-scraping-guide.com, a free public reference on web data collection. Writing that taught me what I would carry into a classroom: the hard part is not explaining what a method does, it is making visible the judgement behind choosing it, because students inherit the confidence of their instruments unless someone shows them where those instruments fail. I would want to teach exactly that, and I expect to need training in university pedagogy.

On dissemination, this project ends in advice to the public and that is much of its appeal. I would want the work to reach the research community through the four papers set out in my proposal, Danish newsrooms through the outlet-level findings, which are directly actionable for editors, and the public through the same plain writing I already do in a technical field. I would release the codebook, the parsing and resolution tooling and the reporting standard openly, because a hub protocol nobody else can run is not a method.

On what I have not done: I have no publications. I have never run an interventional or interview-based study, never recruited participants face to face, and never been the controller for a consent-capture operation, which is the part of the hub I would be learning rather than importing. I have not worked with panel survey data. My proposal is sized to that, the interview and survey work sits with colleagues who own it, and I would rather learn those parts here, from people who do them properly, than claim I already know them.

I have spent seven years turning what people say online into numbers, and the last two becoming uncomfortable with how easily those numbers travel. I would like to spend the next three finding out what they are worth, on a question where the answer reaches people rather than dashboards.

Yours sincerely,

Asad Ikram

Enclosed: research proposal · curriculum vitae · master's degree certificate and transcript · bachelor's transcript · Ikram and Haig manuscript with co-authorship statement