

Asad Ikram

Lahore, Pakistan · +44 7482 376417 · asad.ikram53@gmail.com · linkedin.com/in/asad-ikram98

Associate Professor Kim Andersen

Centre for Journalism, Department of Political Science and Public Management

University of Southern Denmark, Odense

Re: PhD student with a strong computational profile, Navigating Negativity: What is a Healthy News Diet? (SP1)

Dear Associate Professor Andersen,

I found the question behind this application by accident, about eight years ago, while doing work nobody would call research. I was building collectors, and the more sources I collected the more often they disagreed with each other about the same thing.

At Dubizzle Labs I led a three-engineer team running more than five hundred collectors across fifteen countries, in English, Arabic, Urdu and Hindi, feeding a layered warehouse I redesigned into raw, operational and analytical layers. The interesting part was never the scale. It was that the same event looked different depending on where people were talking about it. On a forum, people argued at length and qualified themselves. On social platforms the same sentiment arrived compressed, sarcastic and stripped of context. When we ran language models over both and put the numbers side by side, the comparison looked authoritative and I could not defend it. We could collect almost anything. We could not say what our data measured. I was twenty-four and I assumed this was a gap in my knowledge rather than a gap in the field.

That question is what I applied to Chevening with. My proposal was that governments should have a centralised social analytics capability so that public decisions could be informed by what people actually say rather than by what a communications team believes they say. I was selected from over seventy thousand applicants across a hundred and sixty countries, and I went to Portsmouth to find out whether the idea survived contact with method.

It survived, but not in the shape I expected. My MSc dissertation, supervised by Dr Ella Haig, analysed roughly 279,000 posts on UK economic policy across Twitter, Reddit, YouTube and Quora, with a fine-tuned classifier at macro-F1 0.878 and human validation at Cohen's kappa around 0.78, a figure I note sits close to the tone reliability your department reported in its own comparative coding work. The accuracy was not the finding. The finding was that the same policy event produced structurally different sentiment and emotion profiles on different platforms, which meant the cross-platform aggregate I had built, and the ones my industry ships every day, rested on an assumption nobody had tested: that the classifier measures the same thing everywhere it reads. Dr Haig and I wrote that up as a methods manuscript treating each platform as a survey mode and applying generalizability theory, measurement invariance and score linking to machine-labelled text. In a simulation with known ground truth, an index pooling platform scores at face value reports the wrong sign of the aggregate on roughly a quarter of days. It is complete but unpublished and unsubmitted, it is on English-language material, and it has not been reproduced on Danish data. I attach it with a statement of my contribution so it can be assessed as documented work rather than taken on my word.

Meanwhile the industry side kept teaching me the same lesson at larger scale. At M+C Saatchi Fluency I design and operate the multi-platform social listening pipeline behind campaign intelligence for UK Government departments and for Amazon, Ford, Nike and Reckitt. It is not a scraping job. Sprinklr and Brandwatch feeds and custom collectors land in a governed warehouse on AWS, sixty-eight collectors across ten platforms daily into a 156-column master table, and the analytical work sits on top of it: sentiment scoring, topic and entity tagging, benchmarking brands and narratives against each other, tracking misinformation dynamics and narrative shifts across millions of conversations. Those reports are read by senior government communications teams. What that taught me is how a measurement becomes a decision. A tag on a post becomes a share-of-voice number, becomes a slide, becomes a communications strategy, and by the time it is a strategy nobody remembers that the first step was a model with an error rate.

ArtemisAI is where I stopped consuming that layer and started building it. As founder and CTO I designed the architecture: six fine-tuned encoder classifiers for sentiment, emotion, toxicity, topic, language and intent, served from SageMaker endpoints, trained on a labelling pipeline where a strong model labels a seed set with its reasoning, decision rules are mined from that reasoning into a prompt for a cheaper model, and low-confidence cases route back for review. On top of the classifiers sit the things clients actually asked for: predicting which posts will underperform before they are published, identifying posting windows, drafting posts, and raising crisis alerts when sentiment moves sharply. For the collaboration feature I built an affinity graph and applied network theory to it, which came directly out of reading the literature and asking whether it would hold in production. That is the habit I want to keep: taking something from a paper and finding out whether it survives contact with real data.

And ArtemisAI is where the question I started with came back and found me. Our two annotation tiers agree on 82.9 per cent of sentiment labels and 74.4 per cent of emotion labels, while the toxicity model reports 96.7 per cent average confidence and the dashboard's accuracy figure turned out, when I read our own code, to be inter-model agreement with no criterion behind it. Nobody built that dishonestly. It is what happens when powerful instruments are adopted faster than the discipline to check them. I had built the exact thing I had been suspicious of at twenty-four, and I could not fix it from inside, because in industry there is no budget for stopping to ask what the number means.

That is why I am applying to do this academically, and why to this project specifically. SP1 asks how the amount and character of negativity differ across news outlets, social media platforms and individual news diets. That is my question with a construct I care about attached to it, and with stakes that are not commercial: the project ends in public advice about news diets, and a recommendation about how much negative news is healthy is only as good as the measurement underneath it. Your review with Toff and Ytre-Arne names the blind spot that we know little about the role of the news media and intermediaries in avoidance, and measuring negativity at outlet, platform and feed level is a measurement of exactly that contribution. Your finding with Shehata of a preference for positive news is a reminder that the negativity bias is open rather than premised. And Professor Hopmann's comparative coding work is the reason my proposal codes tone at actor-in-story level and reports Krippendorff's alpha per construct before it reports an F1. My proposal sets out the measurement argument in full; what I want to say here is that it is not an idea I arrived with, it is the thing eight years of building these systems kept forcing on me.

The hub work is the part I can do on day one. Managing donation infrastructure in public spaces is a collection problem with a research consequence, and it fails in two ways I have been paid to prevent for seven years. Download packages change structure without notice, and their folder and key names change with the account's language setting, which is decisive in Denmark where donors' phones will split between Danish and English; the standing defence is golden fixtures per platform and locale, versioned parsers stamped onto every record, and canary donations from project-owned accounts. And donated entries are identifiers and links rather than text, so a negativity classifier cannot read them until every item is resolved back to its content at scale, which Hase and colleagues name as the blocker that makes content analysis of donated feeds unfeasible. I have built the automated repair layer for the first problem and the resolution infrastructure for the second, at Fluency, at Dubizzle, at Fix.com, VendueTech, Prefe and CXG, and across more than three hundred freelance engagements.

On qualifications, I would rather be direct than let a document raise the question. My master's is a UK one-year taught MSc of 180 UK credits, conventionally mapped to 90 ECTS, awarded with Distinction in December 2025 with a dissertation mark of 78 and a credit-weighted average of 70.5. It sits on a four-year, 130-credit-hour bachelor's in computer science, so the total study load exceeds a three-plus-two structure even though the master's alone does not reach 120 ECTS. I read the advert's provision for equivalent documented qualifications as the one this is for, and both transcripts are attached so the judgement rests on documents. I would welcome being told early if it is a barrier.

On teaching, I have none in a university setting and I will not dress up what I have. I led three-engineer teams at Dubizzle Labs and VendueTech, mentored two junior engineers at Prefe, and I write and maintain a free public engineering reference on web data collection that practitioners use. Writing that taught me the thing I would carry into a classroom: the hard part is not explaining what a method does, it is making visible the judgement behind choosing it, because students inherit the confidence of their instruments unless someone shows them where those instruments fail. That is what I would want to teach here. I expect to need training in university pedagogy and would take it.

On dissemination, this project ends in advice to the public, and that is a large part of its appeal. I would want the work to reach the research community through the methods papers, Danish newsrooms through the outlet-level findings, which are directly actionable for editors, and the public through the same plain writing I already do in a technical field. I would release the codebook, the parsing and resolution tooling and the reporting standard openly, because a hub protocol nobody else can run is not a method.

On the working environment, I understand regular physical presence in Odense is expected and I want it rather than tolerate it, having spent seven years working remotely across six countries. I do not speak Danish. My proposal treats that as a design parameter rather than an inconvenience, with annotation, codebook work and the qualitative reading of Danish text done by native speakers while I build the instrument and the infrastructure, and I would begin Danish tuition on arrival. What I have not done, plainly: I have no publications, I have never run a human-subjects study or held ethics approval, I have not conducted interviews, and I have not worked with panel survey data. The proposal is sized to that, the qualitative and survey work sits with colleagues who own it, and the parts I would need to learn I would rather learn here, in a team that does them properly, than claim I already know.

I have spent eight years making numbers out of what people say online, and the last two years becoming uncomfortable with how easily those numbers travel. I would like to spend the next three finding out what they are actually worth, on a question where the answer reaches people rather than dashboards.

Yours sincerely,

Asad Ikram

Enclosed: research proposal · curriculum vitae · master's degree certificate and transcript · bachelor's transcript · Ikram and Haig manuscript with co-authorship statement