

Asad Ikram

Lahore, Pakistan · +44 7482 376417 · asad.ikram53@gmail.com · linkedin.com/in/asad-ikram98

Associate Professor Kim Andersen

Centre for Journalism, Department of Political Science and Public Management

University of Southern Denmark, Odense

Re: PhD student with a strong computational profile, Navigating Negativity: What is a Healthy News Diet? (SP1)

Dear Associate Professor Andersen,

One sentence in the advert is why I am writing: that a crucial task for the PhD student will be to manage the infrastructure of the data donation hubs where individuals donate their social media data in public spaces. Building and keeping alive collection systems that run against changing platforms is what I have done for seven years, and it is the part of SP1 that is hardest to hire for. But infrastructure alone is not a doctorate, and the reason I want this particular post is the question sitting underneath it, which my proposal sets out: SP1 asks how negativity differs across outlets, platforms and individual feeds, and that comparison presumes one instrument reads all three the same way. The field's own evidence says instruments do not travel between genres and languages. Establishing whether this one does, and putting the three scores on a common scale, is the precondition of the project's own question, and it is the thing I am best placed to contribute.

The capability first, since it is the concrete part. At M+C Saatchi Fluency I run a fleet of 68 collectors ingesting ten platforms daily into a governed master table, with an automated repair layer that runs as a named CI workflow: when a source changes shape, the system diagnoses the break, regenerates the extraction logic and tests it against live pages before shipping. Before that I led a three-engineer team at Dubizzle Labs running more than 500 collectors across fifteen countries in four languages. What that work teaches is not scale but failure modes, and two of them decide whether SP1 produces data at all. The first is schema drift: download packages change structure without notice, and their folder and key names change with the account's language setting, which is decisive in Denmark where donors' phones will split between Danish and English. Golden fixtures per platform and per locale, versioned parsers stamped onto every record, and canary donations from project-owned accounts are the standing defence, and the field's published response to this problem is still manual repair. The second is that donated entries are identifiers and links rather than text, so a negativity classifier cannot read them until every item is resolved back to its content at population scale; Hase and colleagues name this as the blocker that makes content analysis of donated feeds unfeasible. That is an engineering problem with a measurement consequence, because resolution fails for deleted and restricted items and biases the recovered feed toward what is still live.

The research side is the same question in a different register, and it is why I stopped being satisfied with building these systems. My MSc dissertation at Portsmouth analysed roughly 279,000 posts on UK economic policy across four platforms, with a fine-tuned classifier at macro-F1 0.878 and human validation at Cohen's kappa around 0.78, a figure I note is close to the tone reliability reported in your department's own comparative coding work. The finding that mattered was not the accuracy. It was that the same policy event produced structurally different sentiment and emotion profiles on different platforms, which meant the cross-platform aggregate rested on an assumption nobody had tested. With my supervisor Dr Ella Haig I wrote that up as a methods manuscript treating each platform as a survey mode and applying generalizability theory, measurement invariance and score linking to machine-labelled text; in a simulation with known ground truth, an index pooling platform scores at face value reports the wrong sign of the aggregate on roughly a quarter of days. I should be precise about its status: it is complete but unpublished and unsubmitted, it is on English-language material, and it has not been reproduced on Danish data. I attach it with a statement of my contribution so it can be assessed as documented work rather than taken on my word. The same problem shows up in my own company, where two annotation tiers agree on 82.9 per cent of sentiment labels and 74.4 per cent of emotion labels while the dashboard reported an accuracy figure that was inter-model agreement with no criterion behind it. I fine-tune six encoder classifiers there, including sentiment and emotion; I have not fine-tuned a generative model, and I would treat an LLM as a first-pass labeller to be validated against a Danish human gold standard, not trusted, which is also what your colleagues' Danish benchmarking work suggests.

On qualifications, I would rather be direct than let a document raise the question. My master's is a UK one-year taught MSc of 180 UK credits, conventionally mapped to 90 ECTS, awarded with Distinction in December 2025 with a dissertation mark of 78 and a credit-weighted average of 70.5. It sits on a four-year, 130-credit-hour bachelor's in computer science, so the total study load exceeds a three-plus-two structure even though the master's alone does not reach 120 ECTS. I read the advert's "or equivalent documented qualifications corresponding to" as the provision this is for, and the transcripts for both degrees are attached so the judgement can be made on documents rather than inference. I would welcome being told early if this is a barrier.

On teaching, I have none in a university setting and I am not going to dress up what I have. What I do have is adjacent and I think relevant: I led three-engineer teams at Dubizzle Labs and VendueTech, mentored two junior engineers at Prefe, and I write and maintain a free public engineering reference on web data collection that practitioners use. Writing that taught me the thing I

would carry into a classroom, which is that the hard part is not explaining what a method does but making visible the judgement behind choosing it, and that students reproduce the confidence of their instruments unless they are shown where those instruments fail. If I taught within this centre I would want to teach exactly that: how to read an automated measure sceptically. I expect to need training in university pedagogy and would take it.

On dissemination and practical translation, this project ends in advice to the public about news diets, and that is a large part of its appeal. A recommendation about how much negative news is healthy is only as good as the measurement underneath it, and a number that is not comparable across the sources composing a person's diet cannot support public advice, however carefully it is phrased. I would want the work to reach three audiences: the research community through the methods papers described in the proposal, Danish newsrooms through the outlet-level findings, which are directly actionable for editors, and the public through the same plain writing I already do in a technical field. I would also expect to release the codebook, the parsing and resolution tooling and the reporting standard openly, because a hub protocol nobody else can run is not a method.

On the working environment: I understand that regular physical presence in Odense is expected and I want that rather than tolerate it, having spent seven years working remotely across six countries. I do not speak Danish. In the proposal I treat that as a design parameter rather than an inconvenience, with annotation, codebook work and the qualitative reading of Danish text done by native speakers while I build the instrument and the infrastructure, and I would start Danish tuition on arrival. What I have not done, plainly: I have no publications, I have never run a human-subjects study or held ethics approval, I have not conducted interviews, and I have not worked with panel survey data. The proposal is sized to that, the qualitative and survey work sits with colleagues who own it, and the parts I would need to learn I would rather learn here, in a team that does them properly, than claim I already know.

Thank you for reading. I would welcome the chance to discuss the measurement argument, or the hub, in more detail.

Yours sincerely,

Asad Ikram

Enclosed: research proposal · curriculum vitae · master's degree certificate and transcript · bachelor's transcript · Ikram & Haig manuscript with co-authorship statement